

# Course: Domain Driven Design & Microservices for Architects

Section: SAGA Pattern

<http://acloudfan.com/>

Pragmatic Paths Inc © 2021

Contact: [raj@acloudfan.com](mailto:raj@acloudfan.com)

Discount Link to course:

<https://www.udemy.com/course/domain-driven-design-and-microservices/?referralCode=C5DCD3C4CC0F0298EC1A>

# Distributed Transactions : SAGA

---

Distributed System Transactions

2 Phase Commit **CANNOT** be used

SAGA pattern is used for Distributed Transactions

SAGA are complicated !!!

Multiple frameworks available for realization

Framework hides complexity

1

SAGA pattern

2

Types of SAGA (Orchestration Vs. Choreography)

3

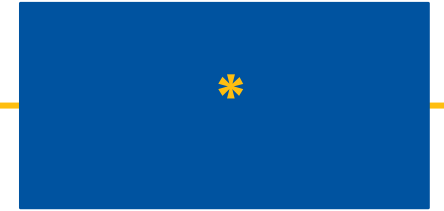
Implementation considerations

4

Case Study : Working with SAGA for ACME Sales booking transaction  
(UML, JAVA, Kafka, PostgreSQL, MongoDB)

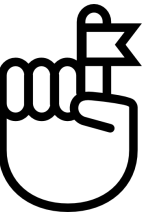
# Data consistency in Distributed Tx

Distributed transactions and data consistency



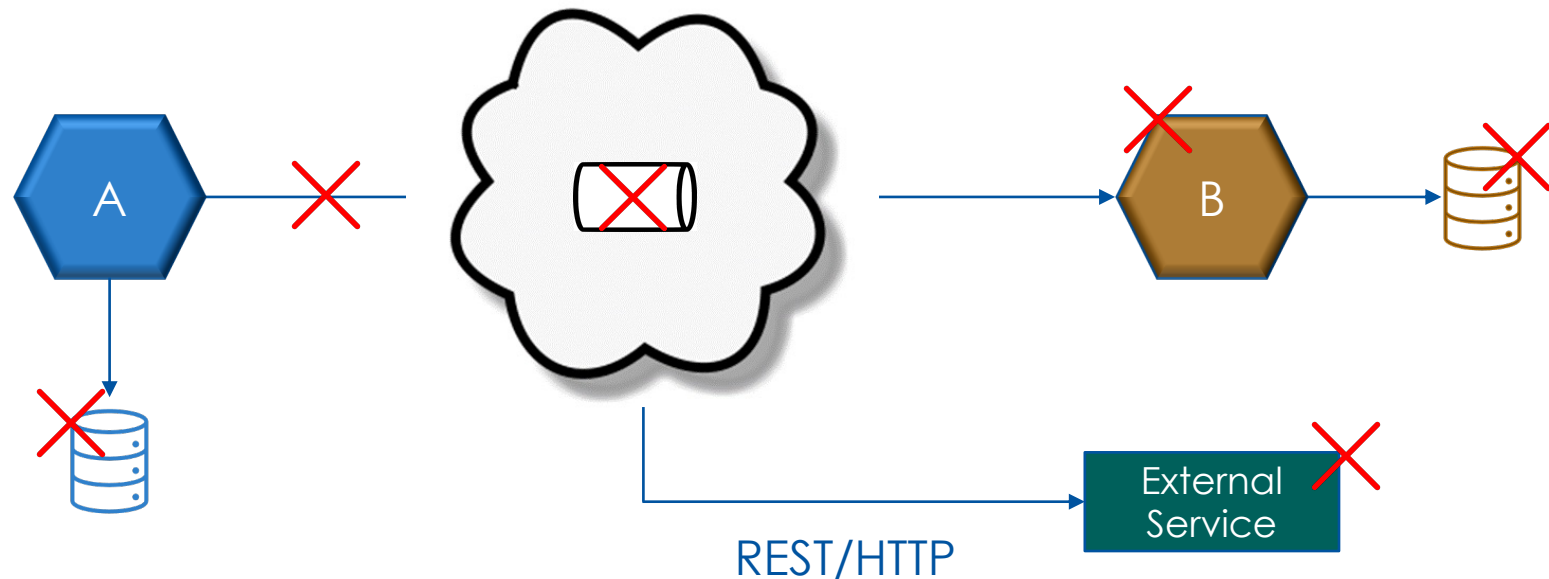
- 1 Data consistency in microservices
- 2 Managing data consistency
- 3 Introduction to SAGA pattern

# Design for failure



Always anticipate that there will be FAILURES

- Identify the "Failure Points" in your architecture

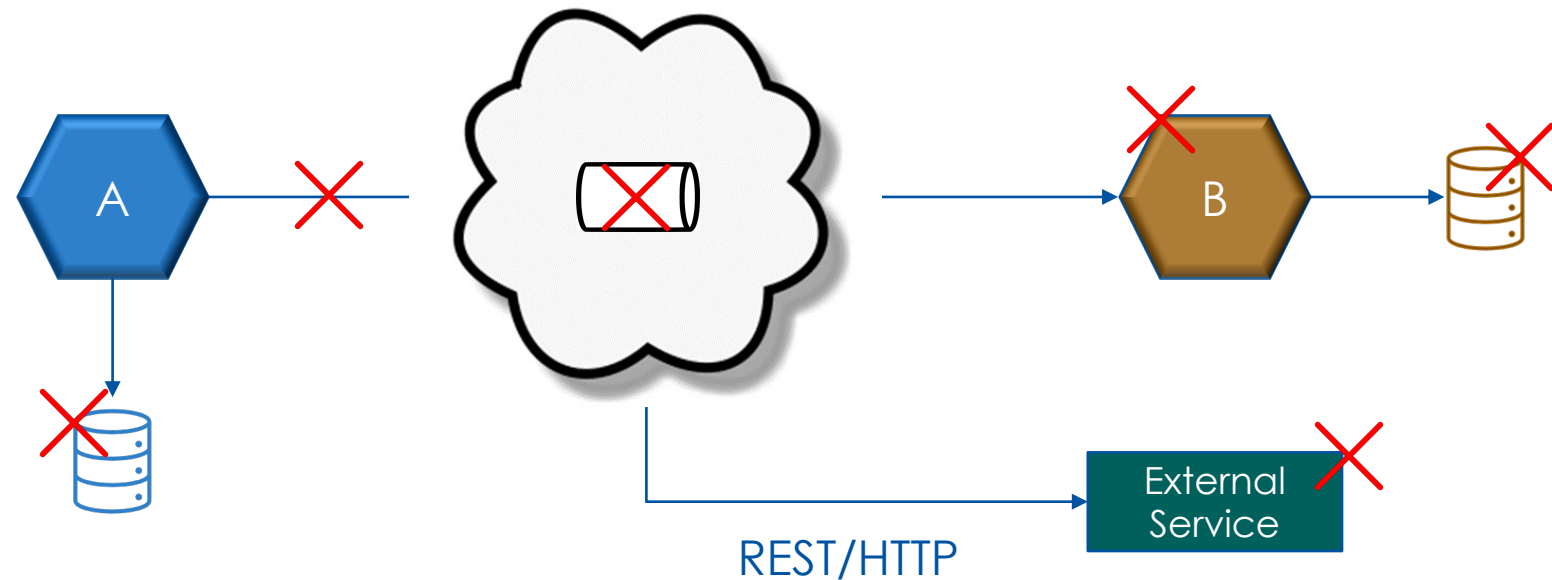


Pro-actively address the "Failure Points"

## Design for failure

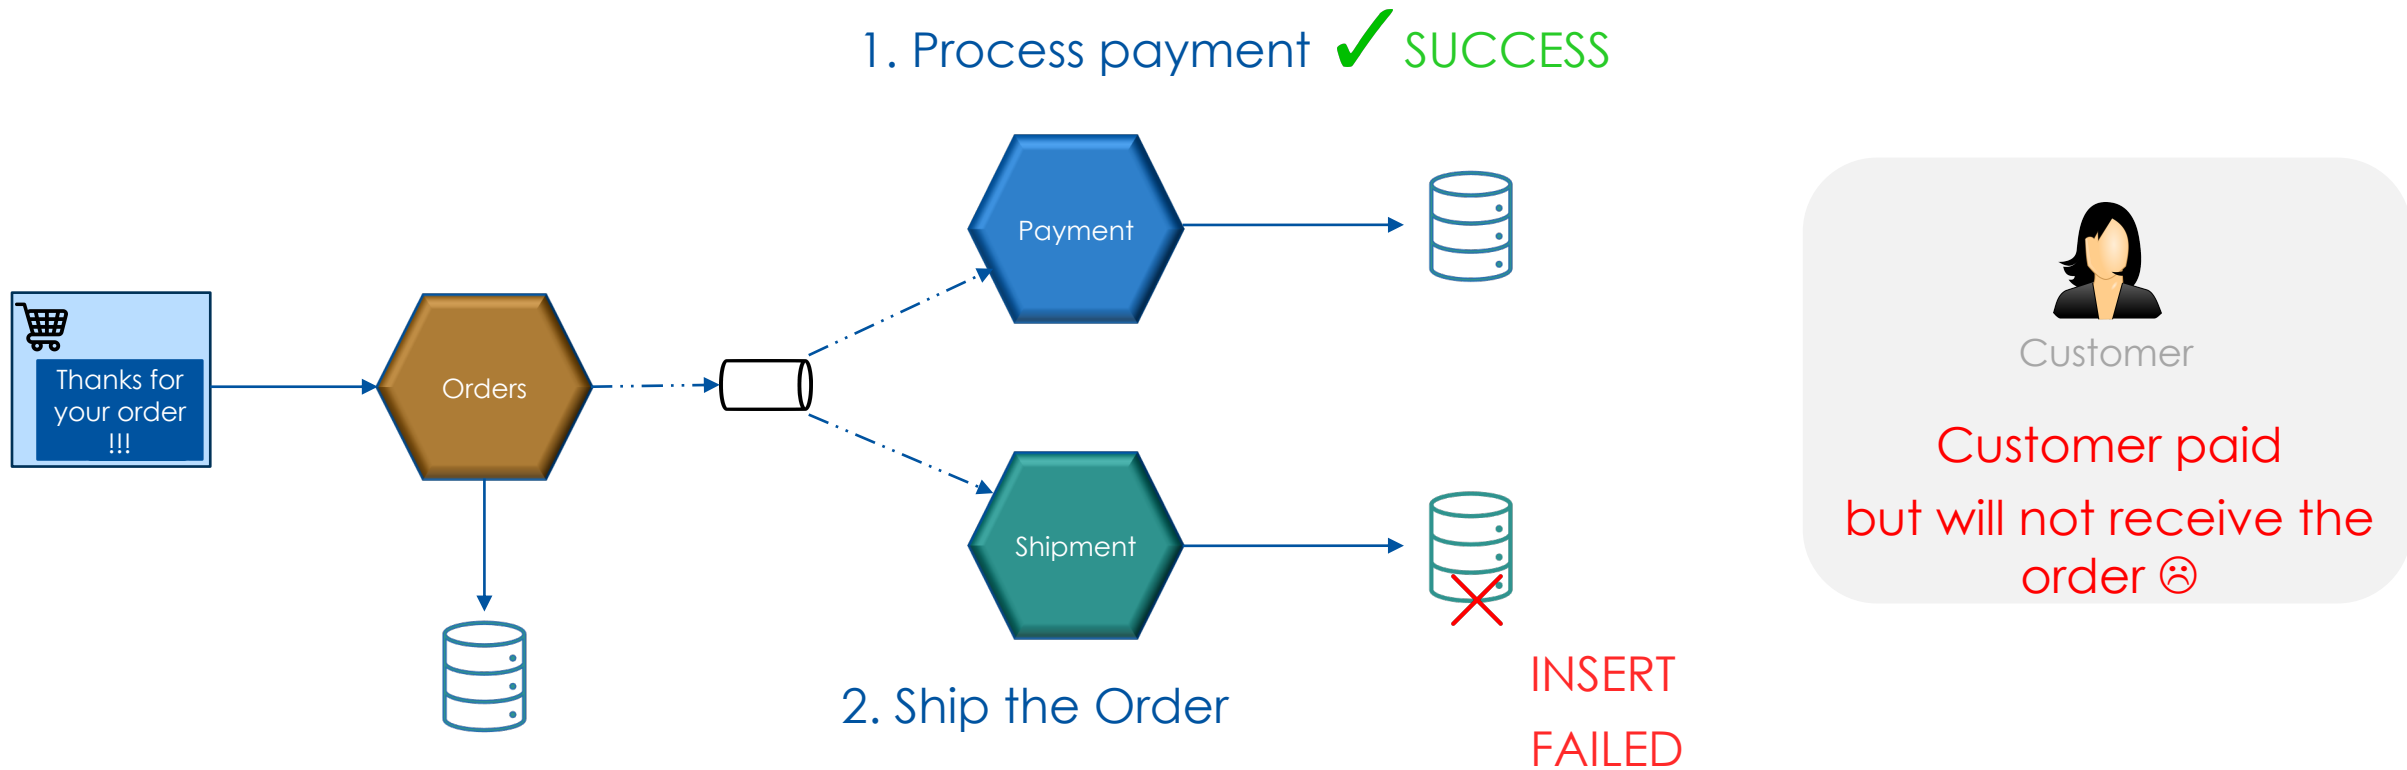


NOT addressing the failures = Inconsistent state of data



NOT addressing the failures = In consistent state of data

## Example: Data inconsistency

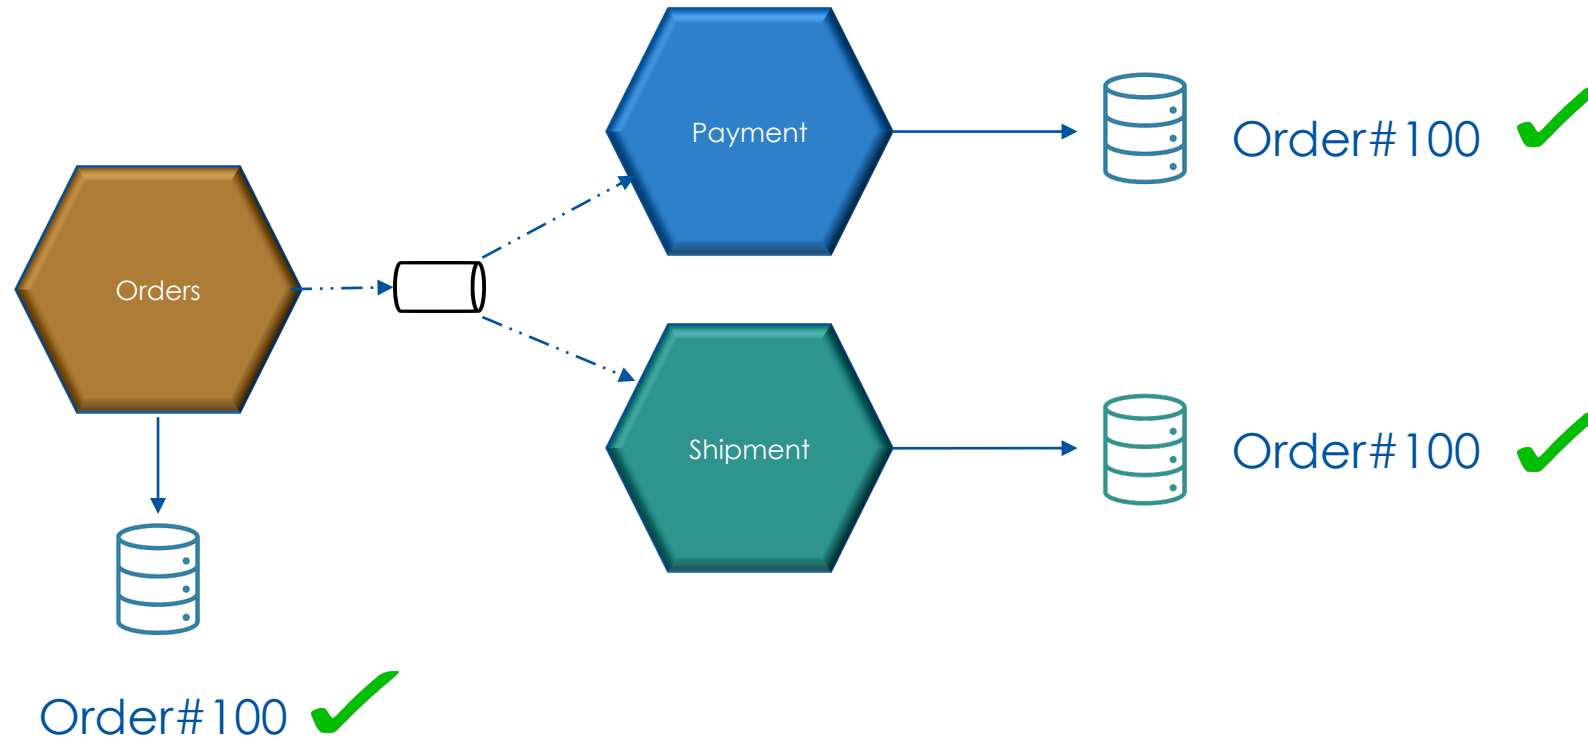


Shipment database is not in sync with other databases !!!



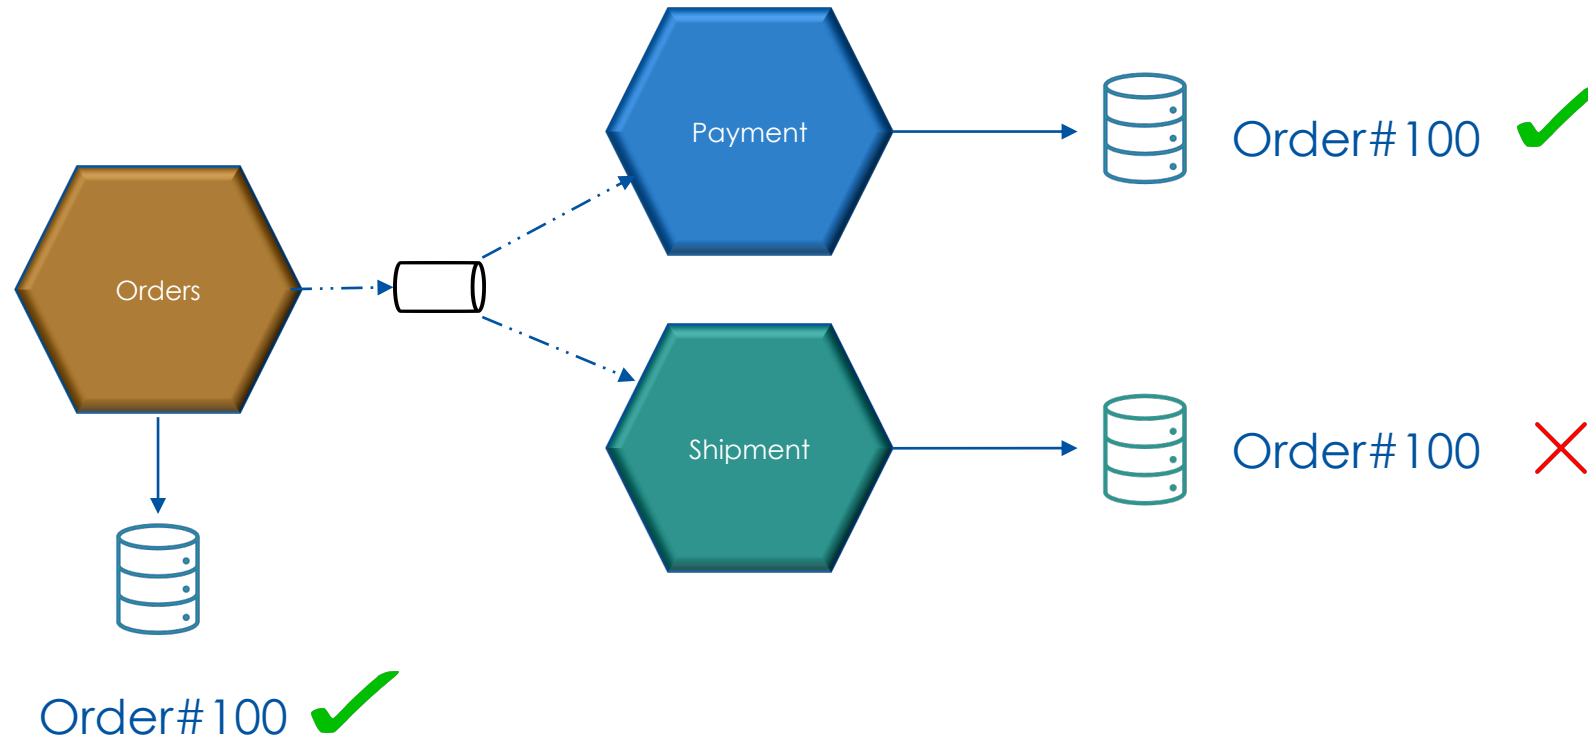
# Objective

Data **MUST** be consistent across all distributed services



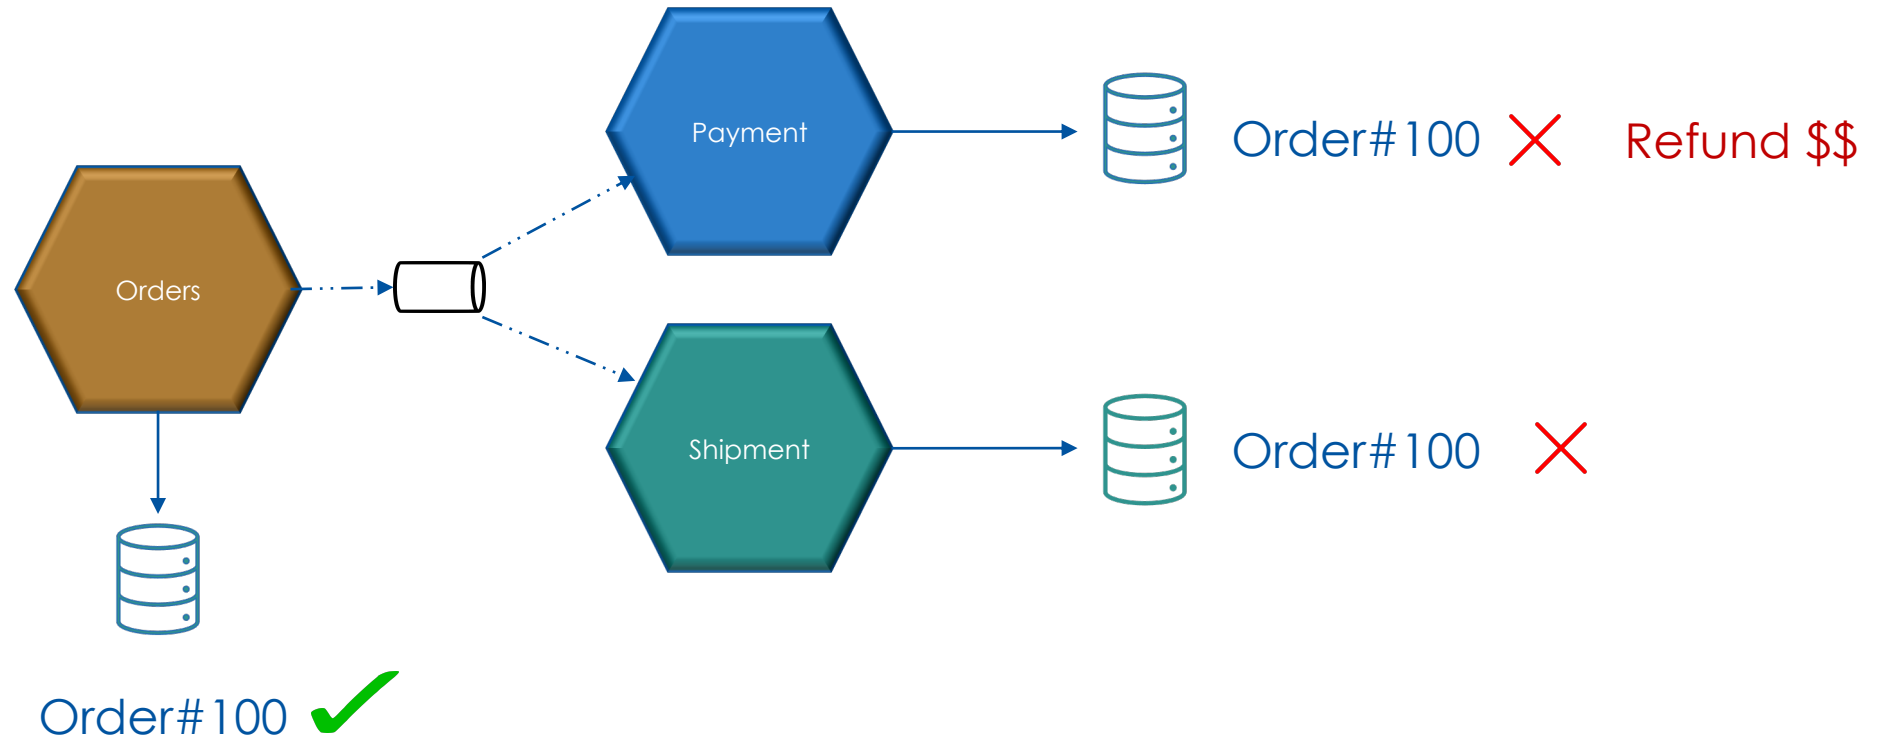
# Objective

Data **MUST** be consistent across all distributed services



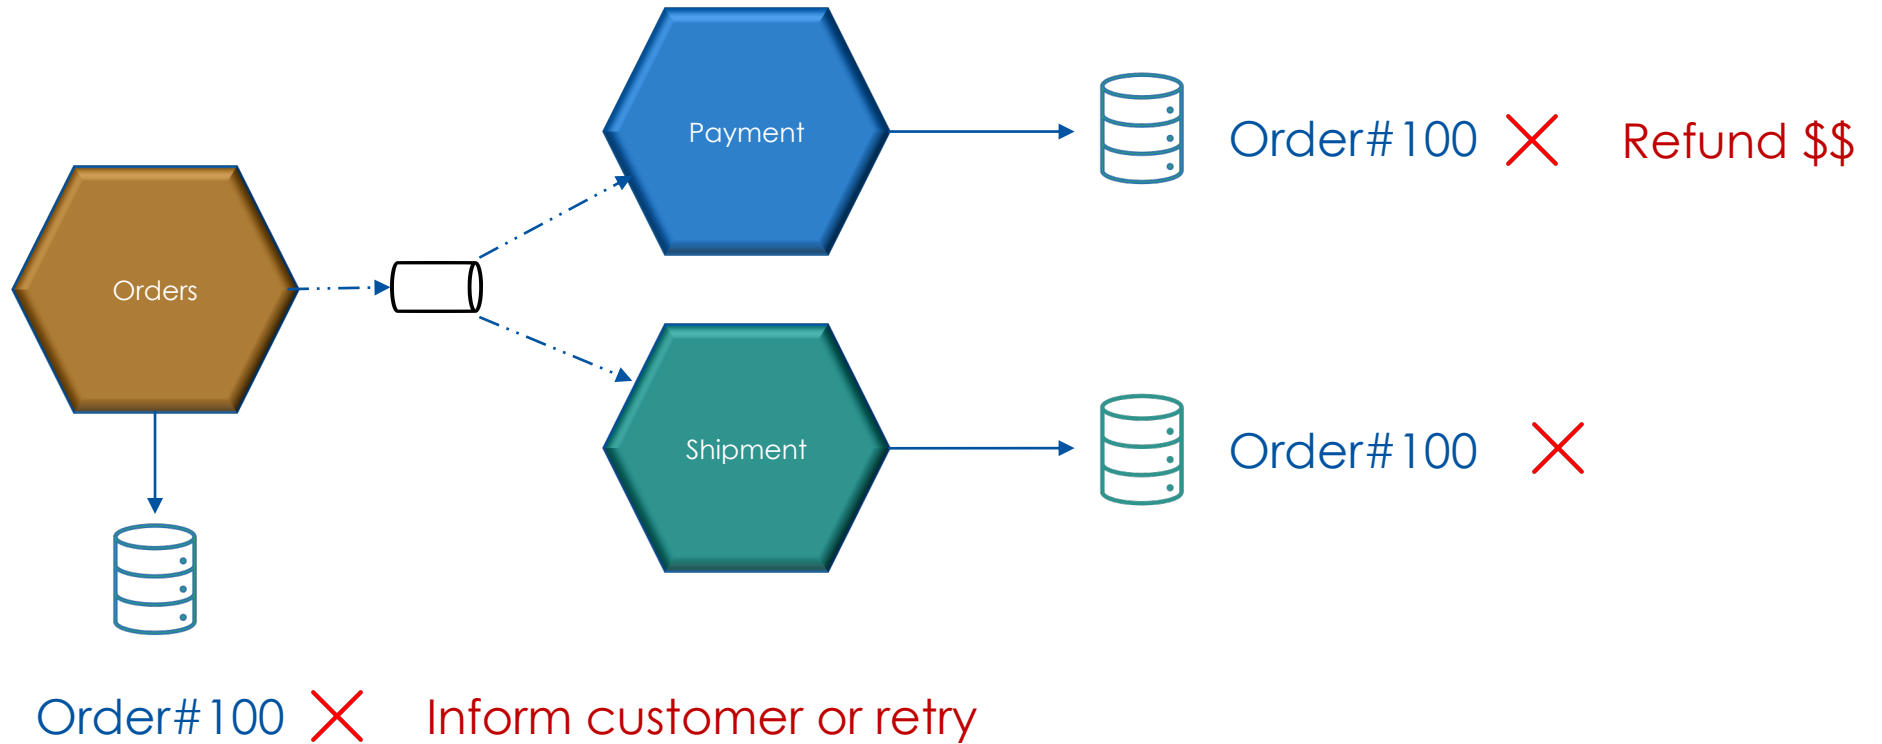
# Objective

Data **MUST** be consistent across all distributed services



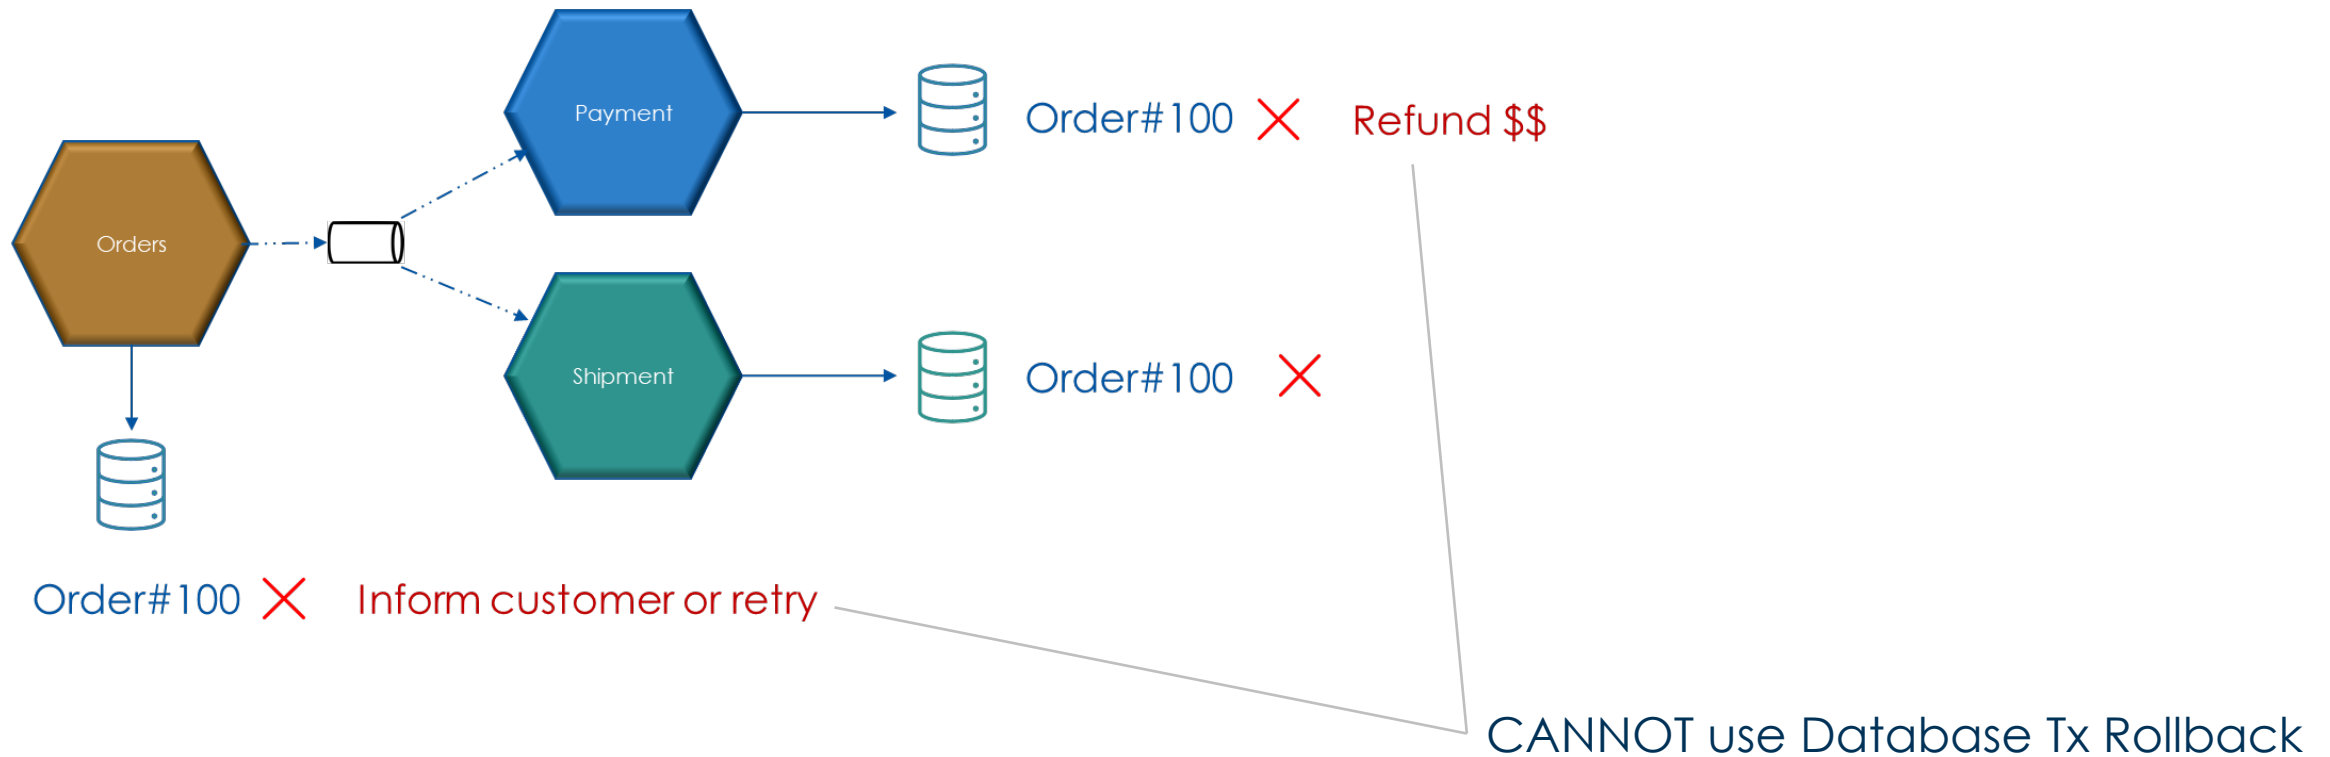
# Objective

Data MUST be consistent across all distributed services



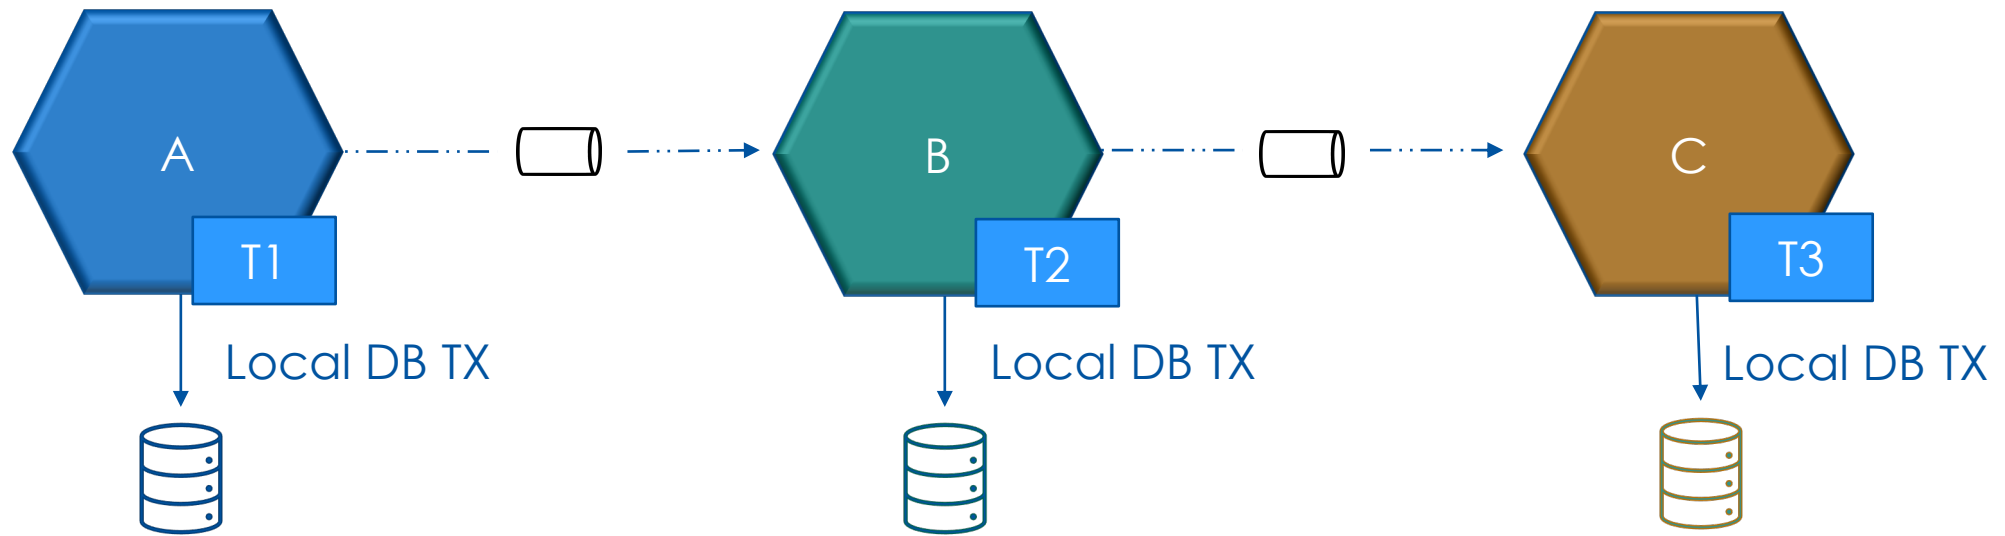
## Local Transactions

**CANNOT** be used for reverting the DB changes



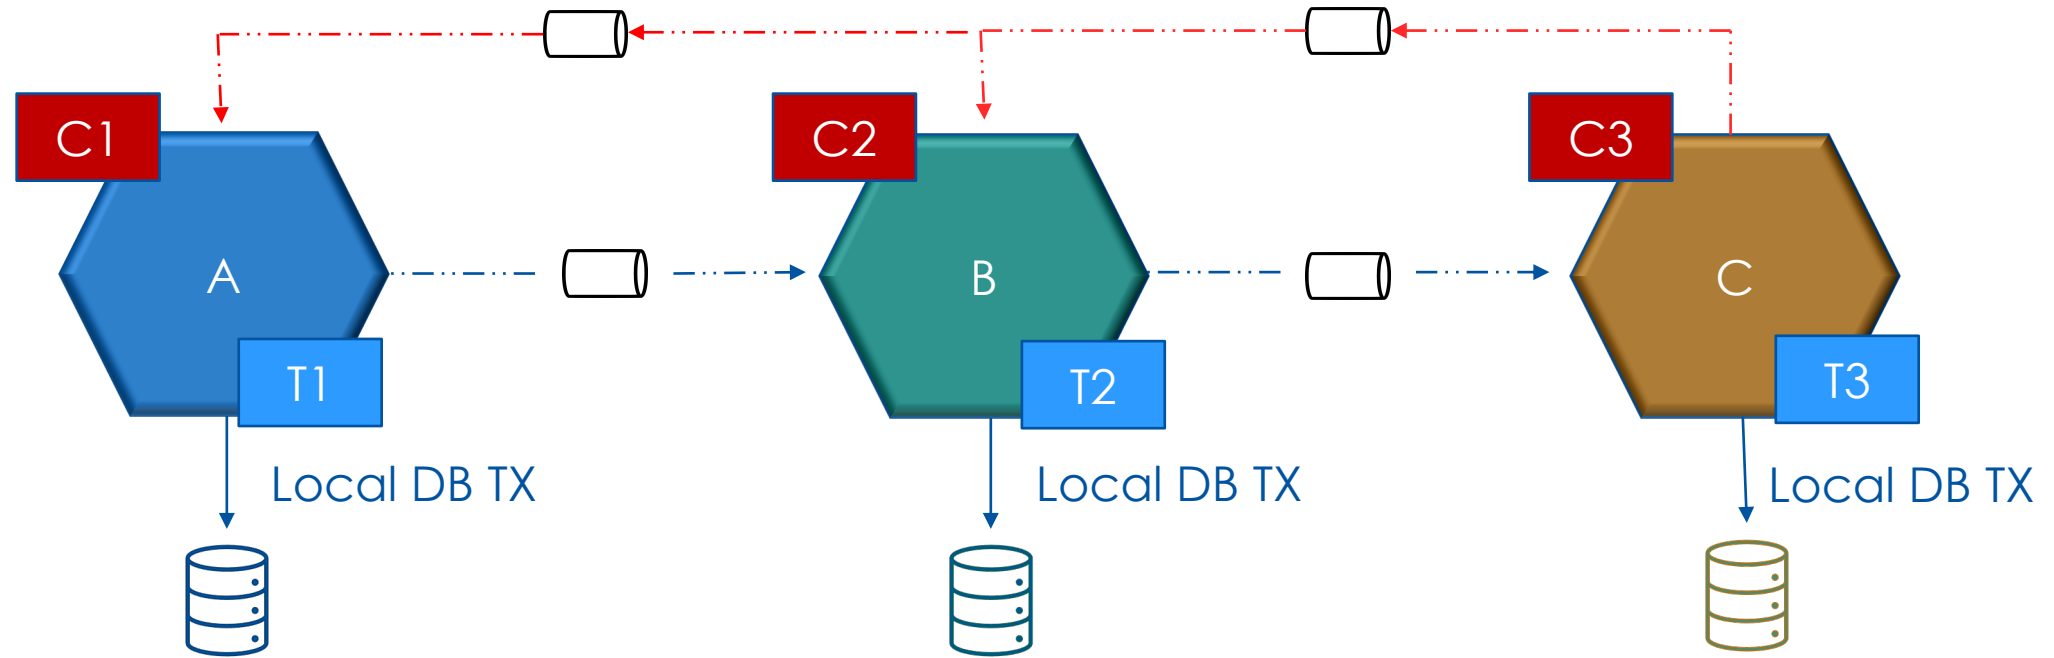
## SAGA pattern

Use Local Transactions coupled with compensating transactions



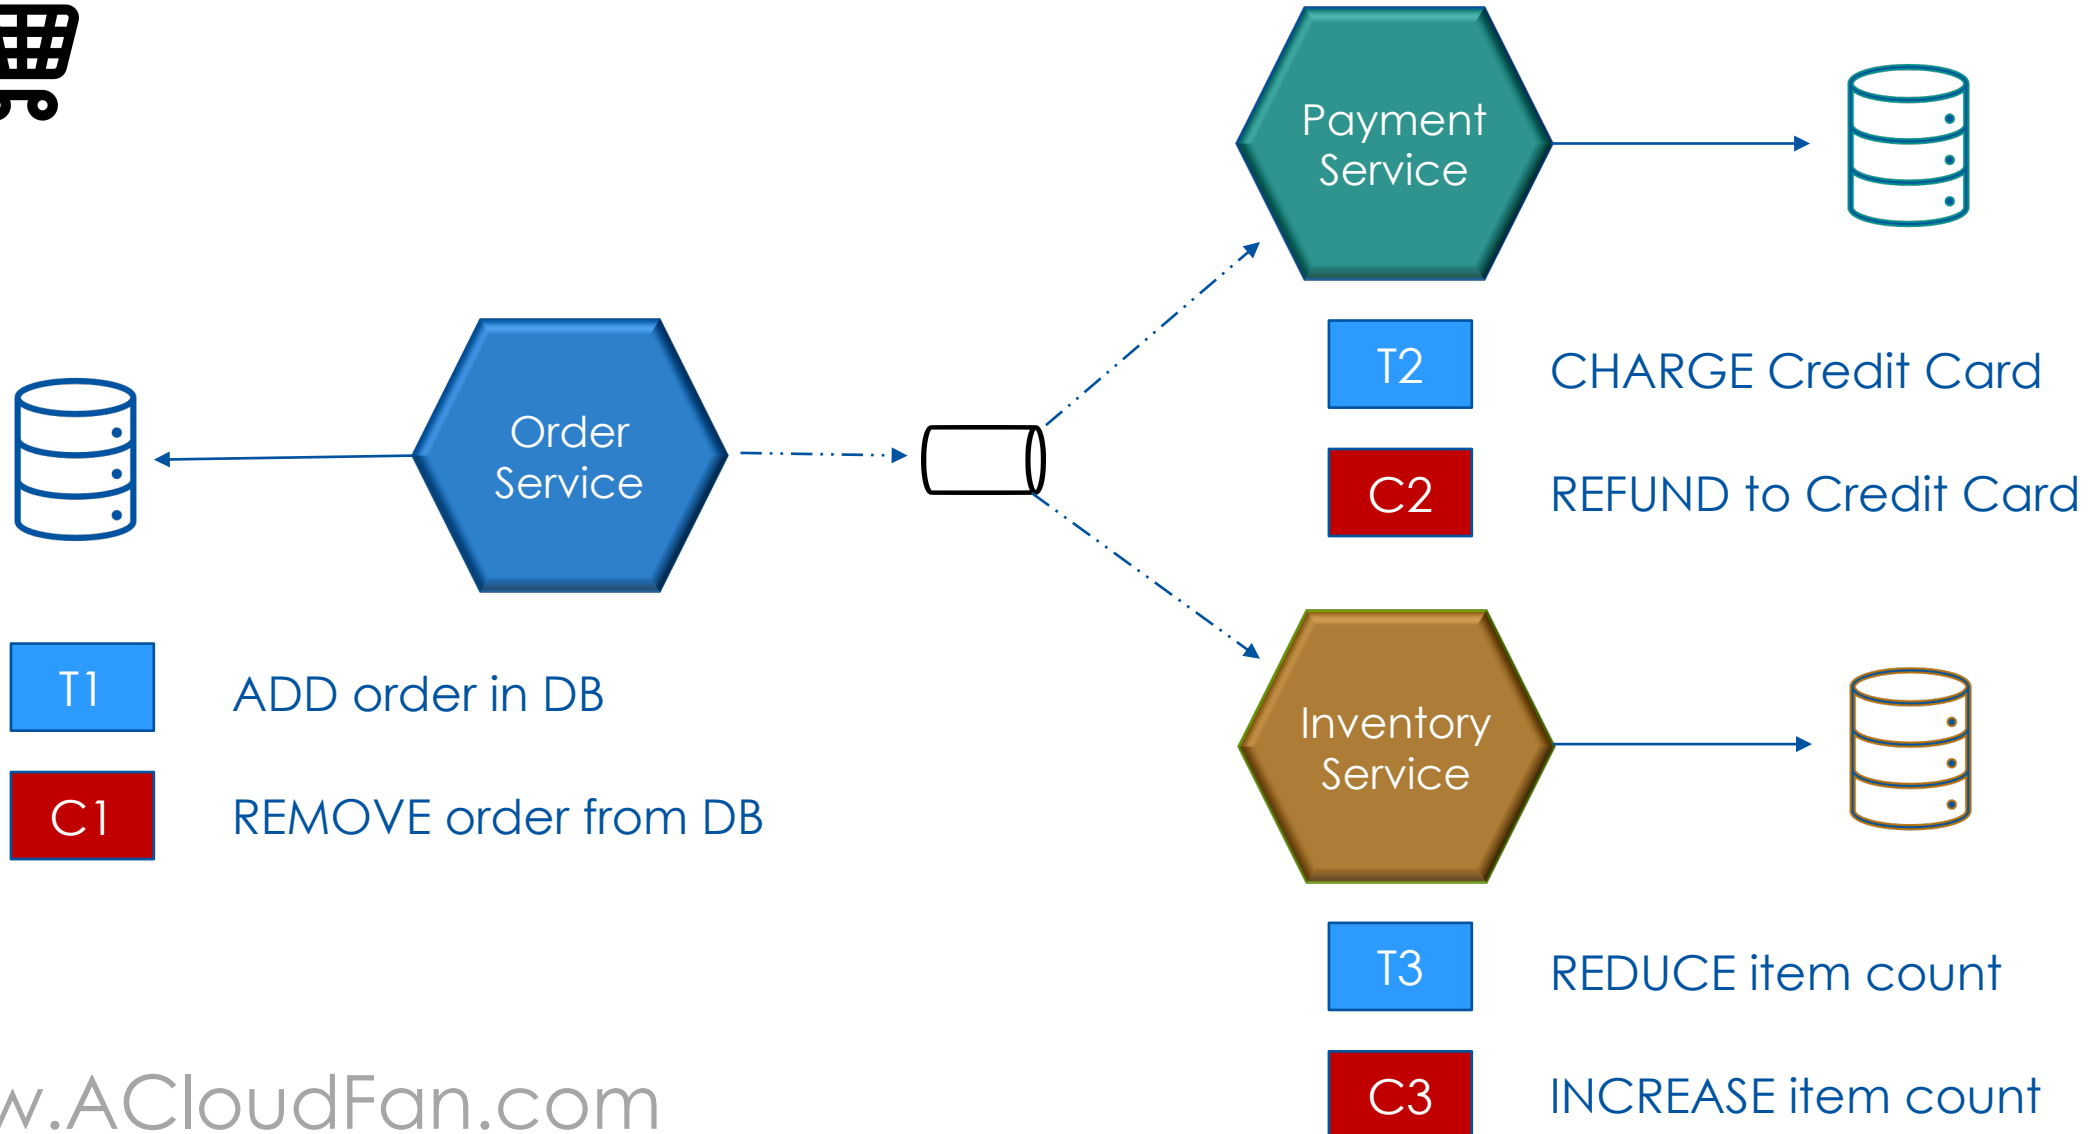
# SAGA Pattern

# Compensating transactions revert the database changes



Requires each Tx to have corresponding Compensating Tx

## Example: Compensating transactions





## History

Introduced in a paper published in 1987 !!

# Sagas

Hector Garcia-Molina, Kenneth Salem  
Princeton University 1987

<https://www.cs.cornell.edu/andru/cs711/2002fa/reading/sagas.pdf>

## SAGA pattern

May be applied to monolithic & distributed systems

- Distributed SAGA = Services are distributed



## Quick Review

SAGA pattern - Managing data consistency across microservices

Uses Local Transactions for persisting Txns to the database

Uses compensating transactions for reverting database changes

# SAGA Call Flow

De-Centralized versus Centralized



# Flavors of SAGA

SAGA may be implemented in two ways



## Event/Choreography

- NO Central component to manage transactions

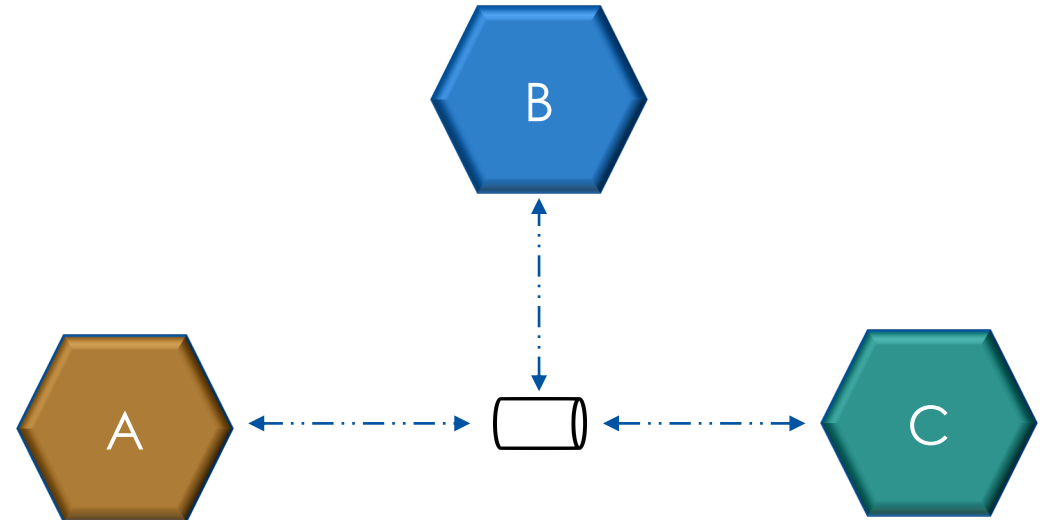
## Command/Orchestration

- Central component for managing the flow
- Central component = SAGA Execution Coordinator (SEC)

## Event/Choreography SAGA

NO central component to manage the SAGA

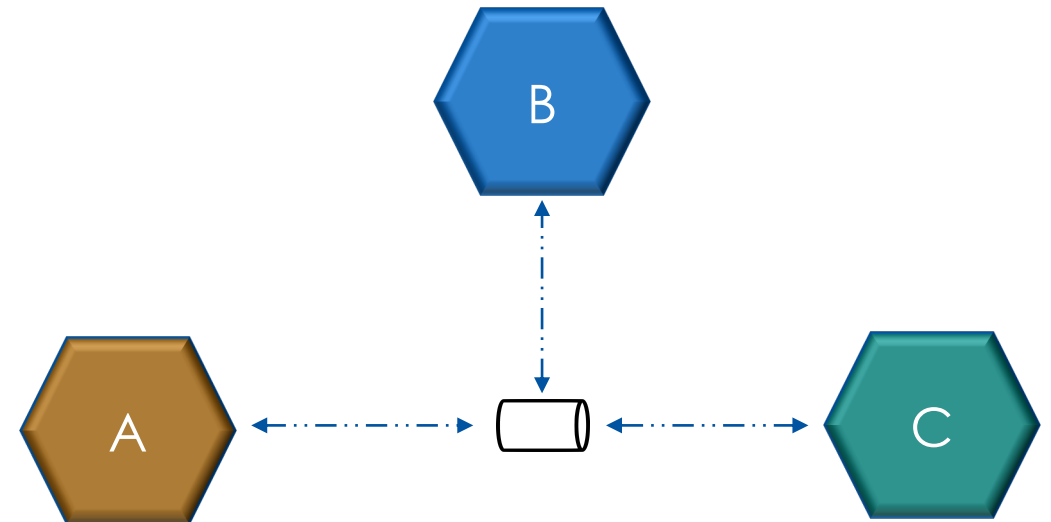
- Services emit & receive domain events
- Services decide on actions independently



# Challenges with Event Choreography

Leads to HIGHLY decoupled services but there are challenges

- Difficult to implement, test and debug
- Out of sequence events
- Coordinating the failure scenarios
- Cyclic dependencies



## Command/Orchestration

Central component manages the calls to services in SAGA



### 1. Domain Object

- Part of the domain model

### 2. Dedicated orchestrator

- Generic i.e., outside of the domain model

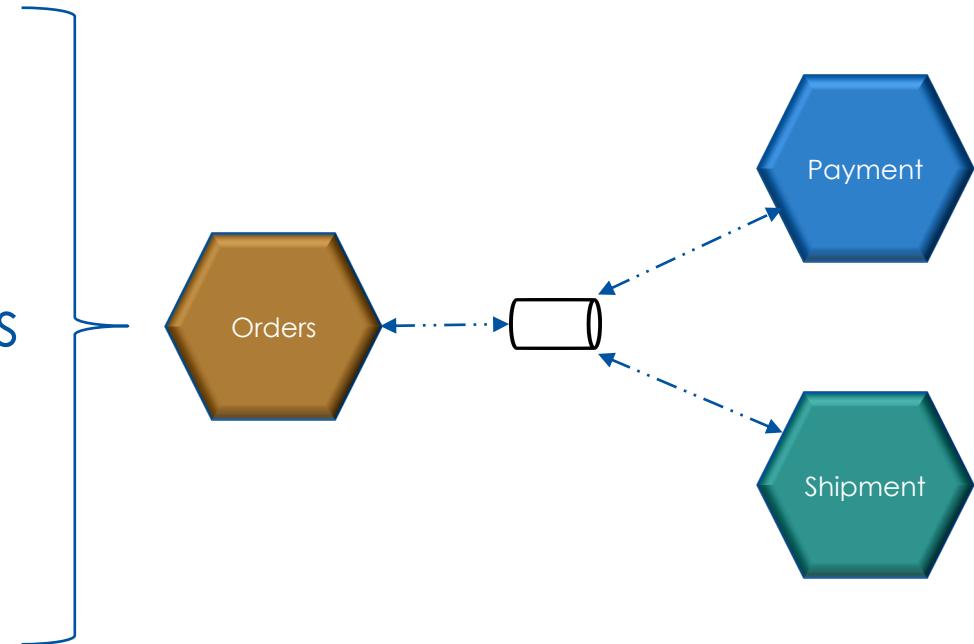
Central Component = SAGA Execution Coordinator



# SAGA Command/Orchestrator

## 1. Key domain object plays the role of SEC

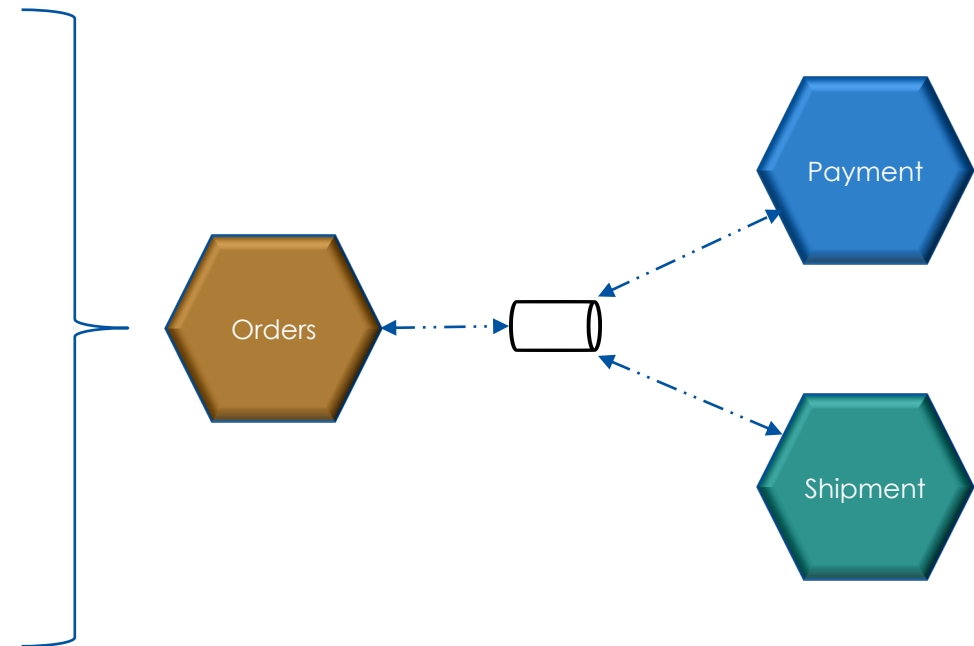
- Manages state of orders
- Initiates the transactions on services
- Manages failures



# SAGA Command/Orchestrator

## 1. Key domain object plays the role of SEC

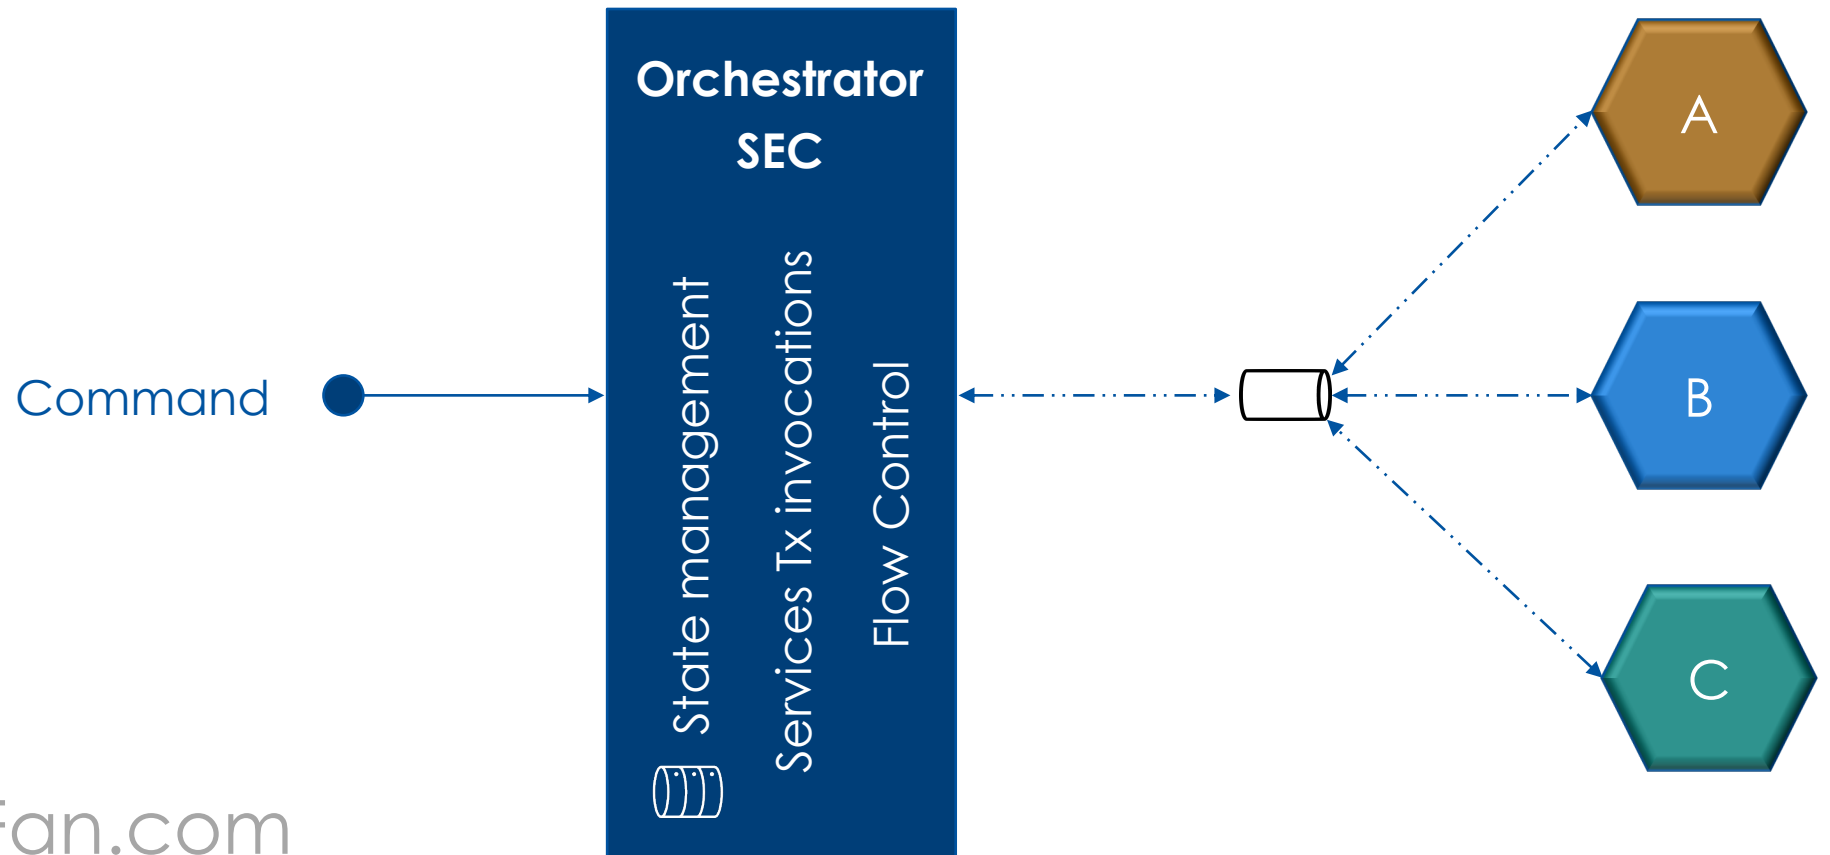
- Uses Command/Reply pattern
- Reliable messaging pattern



State Management will introduce COMPLEXITY in the domain object !!!

# SAGA Command/Orchestrator

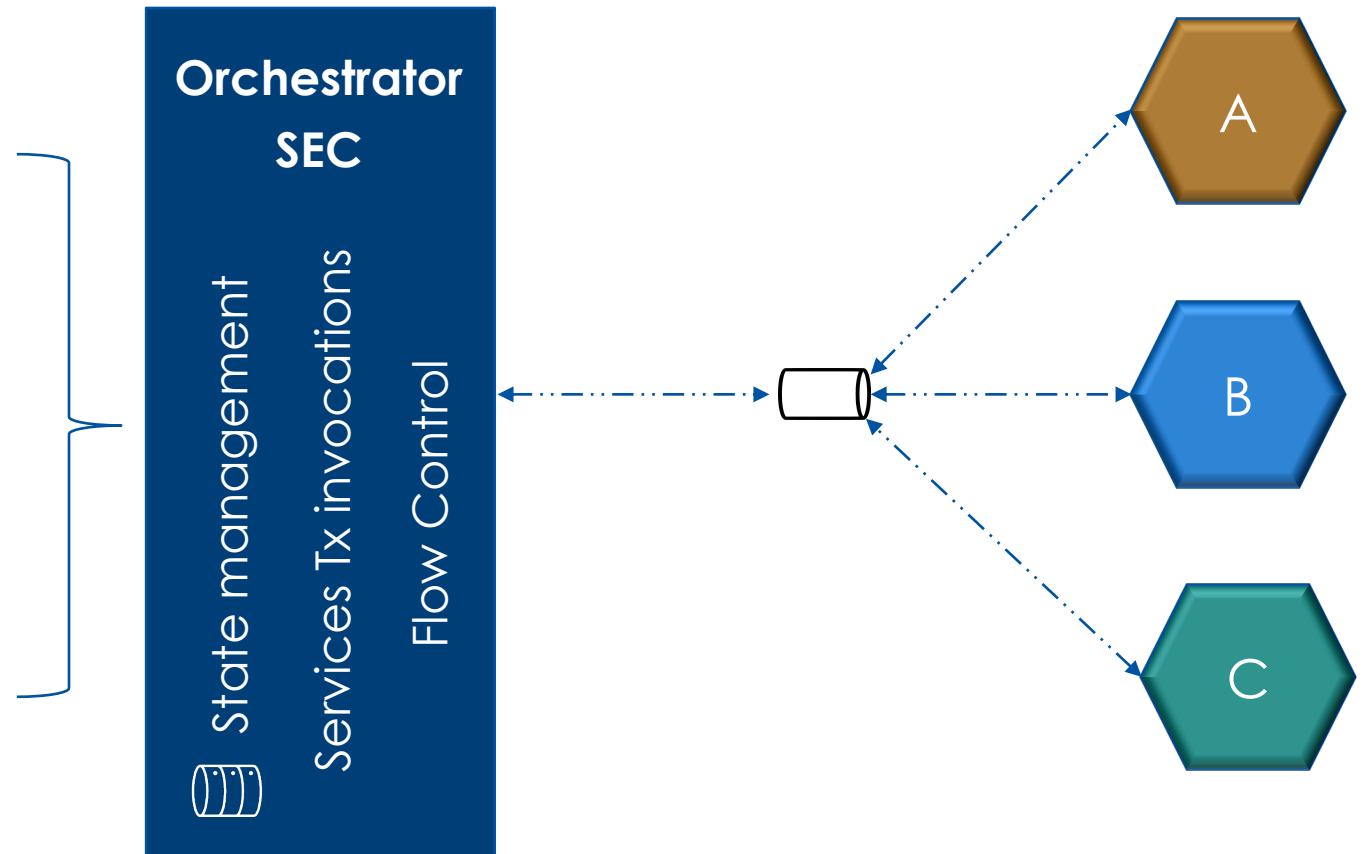
## 2. A dedicated SEC outside of the domain model



# SAGA Command/Orchestrator

## 2. A dedicated SEC outside of the domain model

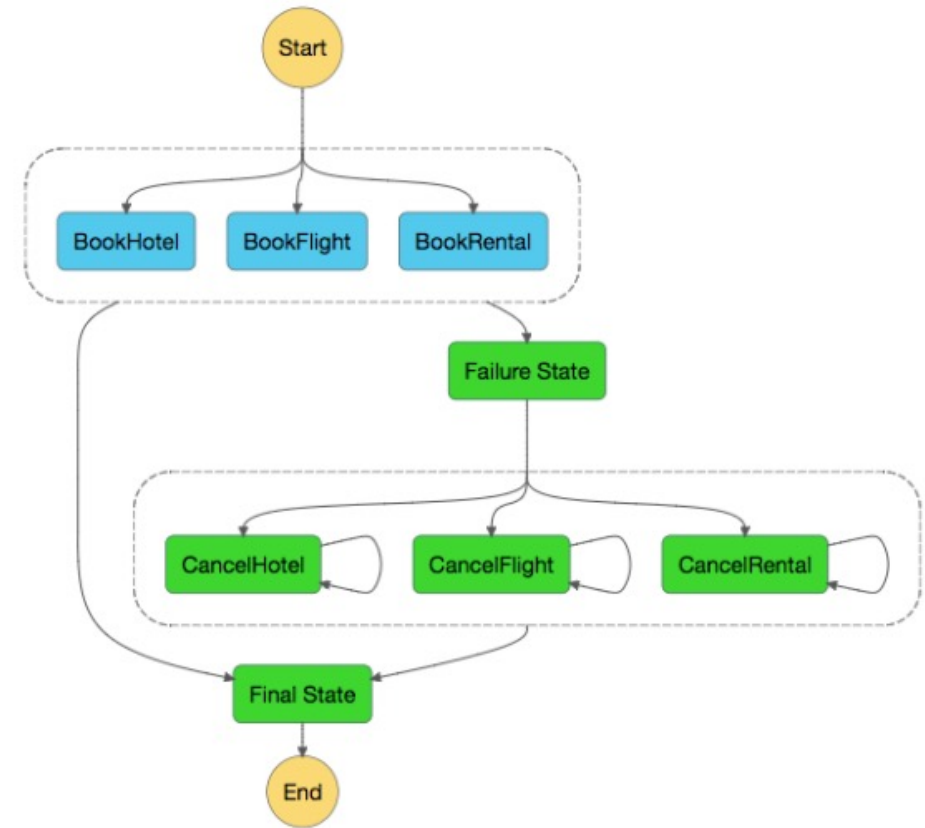
- E.g., use of BPM tool
- E.g., AWS step function
- E.g., Spring framework



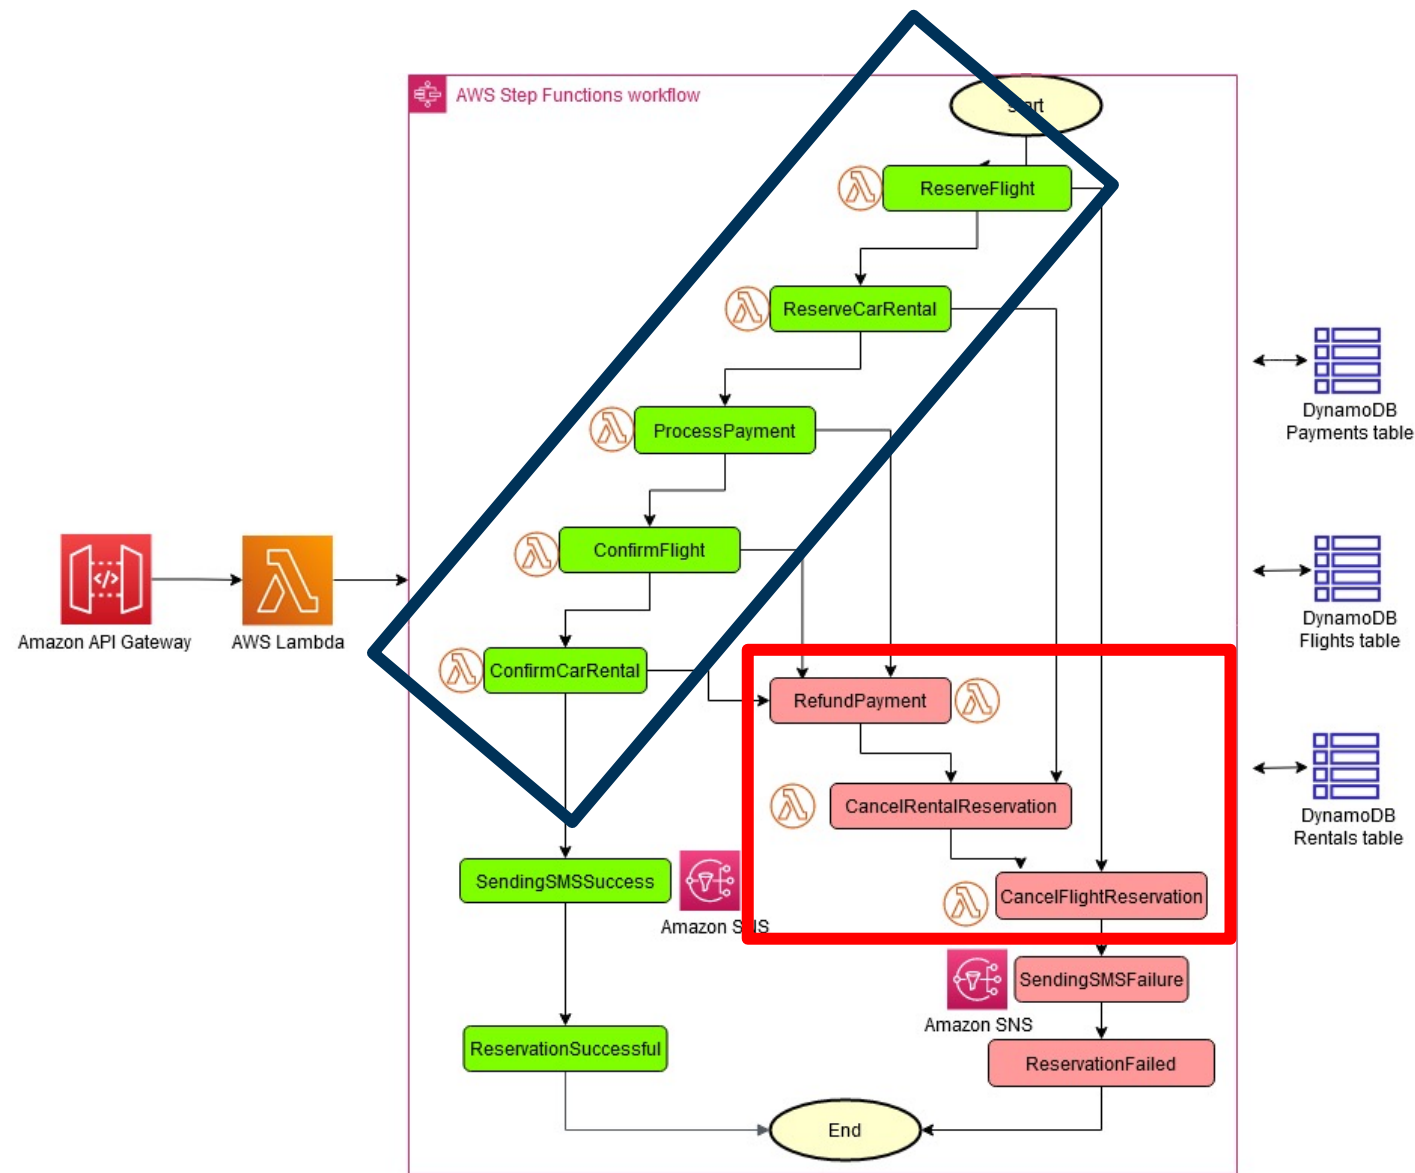
## Example SEC: AWS Steps functions

Define the business process

- Declarative State machine in JSON format
- Business logic "AWS Lambda functions"
- "Step Function" coordinates the execution



# Example: AWS Steps function workflow



## Benefits : Command/Orchestrator

Less decoupled compared & introduces Single point of failure BUT

- Simplicity
- Easier to implement, test and manage
- Rollbacks are easier to manage
- Centralized way to check out the state



## Quick Review

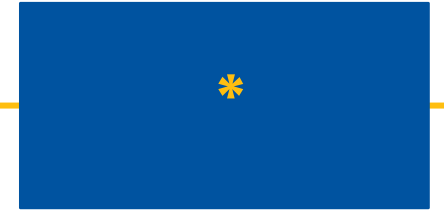
### SAGA

- Event/Choreography = NO central coordinator
- Command/Orchestrator= SAGA Execution Coordinator
  - Domain object
  - External component



# SAGA Design Considerations

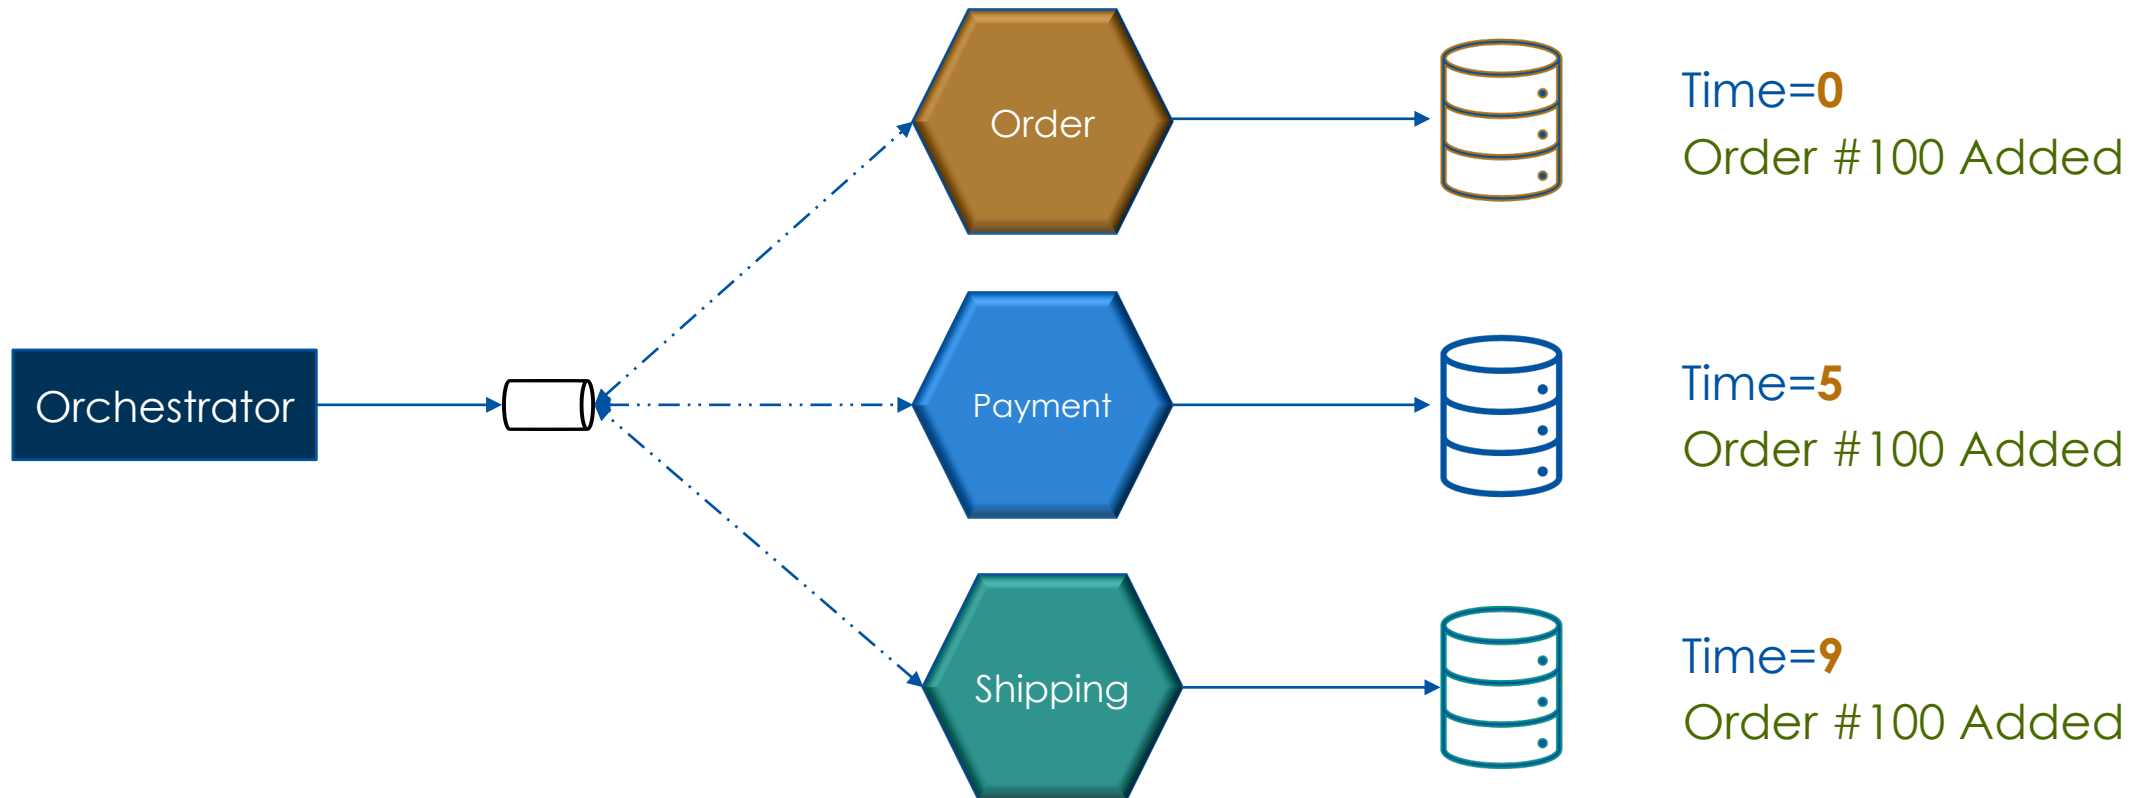
Managing data consistency in Microservices



- 1 Data consistency challenge
- 2 Pros of shared database
- 3 Cons of shared database

# Eventually consistent

Microservices may reflect different states while the Tx is in progress

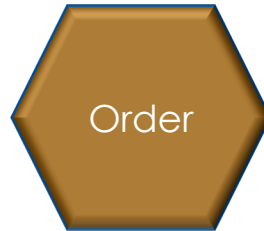


# Isolation Level

SAGA provides isolation level = "READ Uncommitted"

Time=0

Order #100 Added



Time=1

Order #100 processing



Time=..

*pending*



# Isolation Level

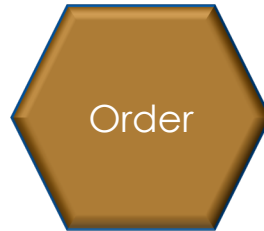
SAGA provides isolation level = "READ Uncommitted"

Time=0

Rollback Order# 100

Time=0

Order #100 Added



Time=2

Order #100 failed

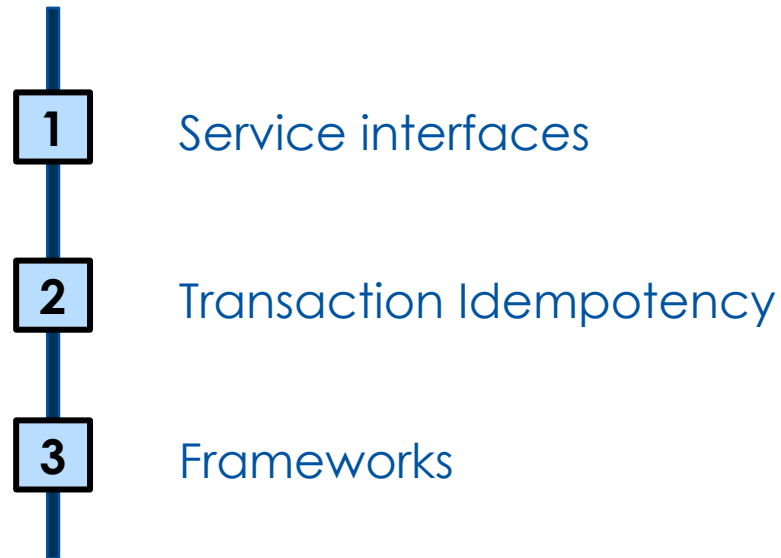
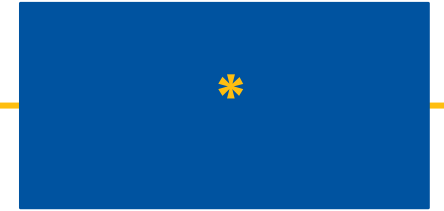
Time=..

*pending*

Data READ is not guaranteed !!

# Implementation considerations

Distributed SAGA implementation approach



## Frameworks

Consider using Frameworks



<https://eventuate.io/>



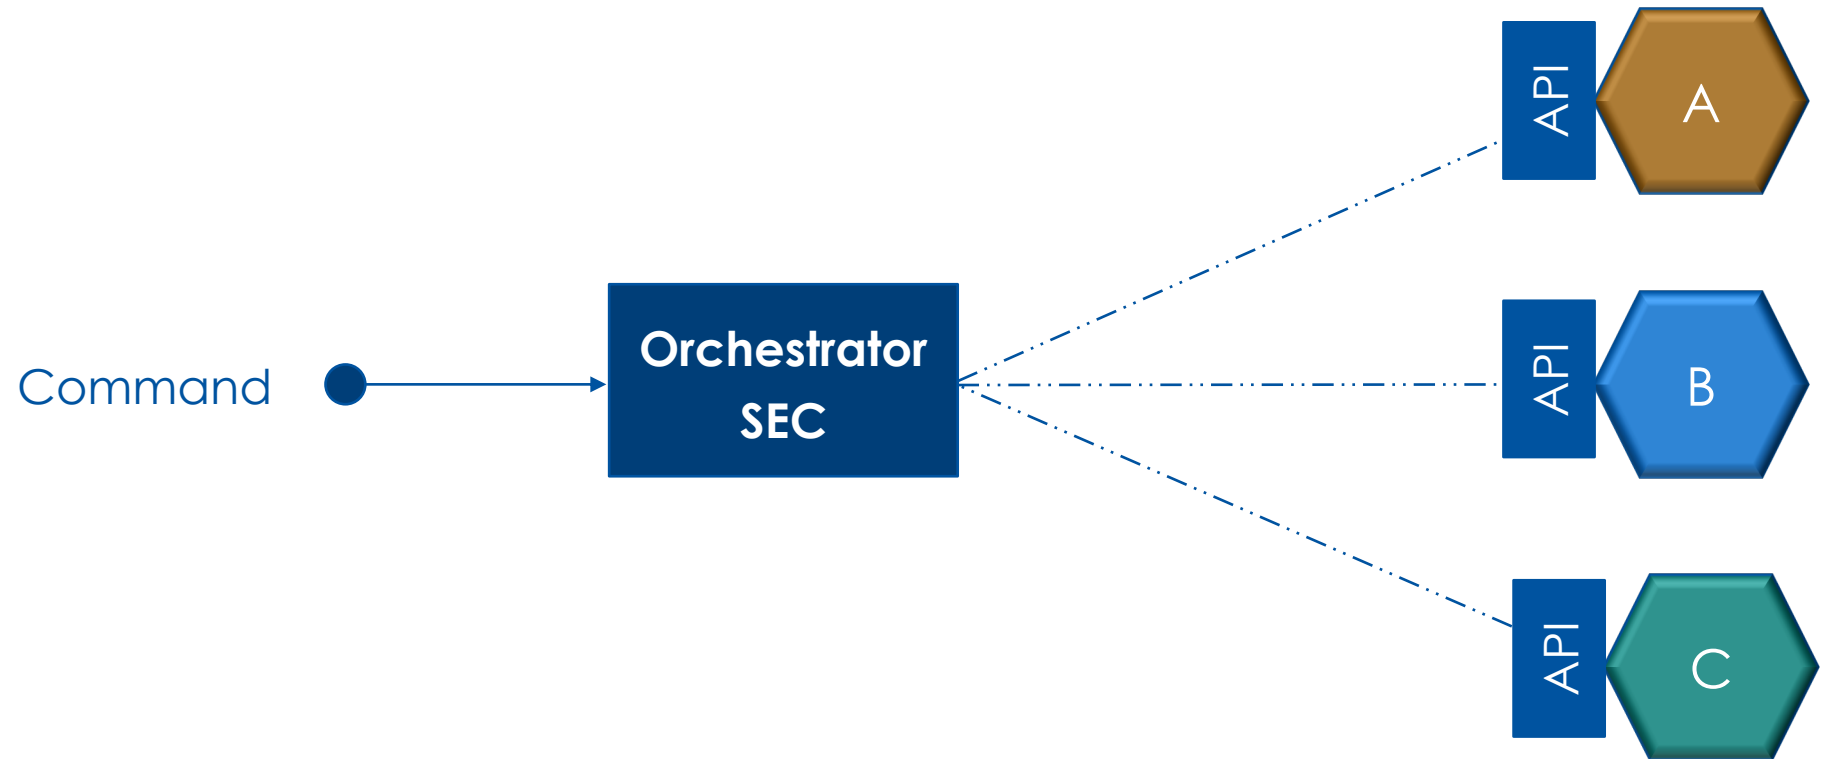
<https://axoniq.io/>



<http://seata.io>

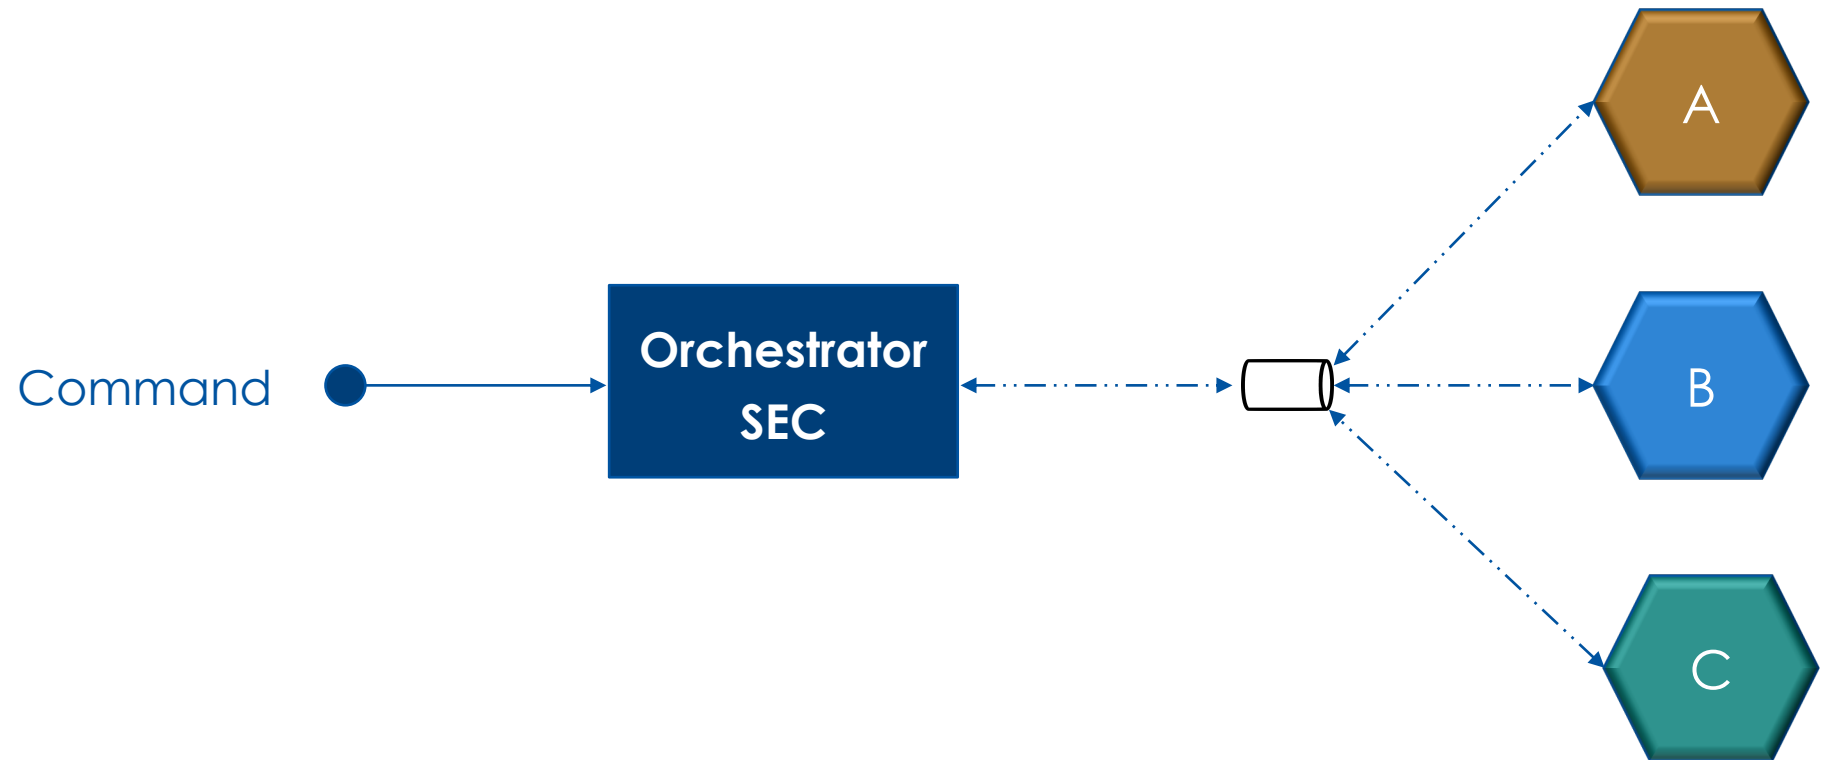
## Synchronous Vs Asynchronous calls

Synchronous calls may be used but messaging is preferred



## Synchronous Vs Asynchronous calls

Synchronous calls may be used but messaging is preferred







\*\*\*OR\*\*\*



# What would you use for implementing SAGA?

## Unique ID per transaction

Each transaction | event has a UNIQUE identity

- Helps with Event Sourcing
- Consumers can identify duplicate transactions

# Idempotent

Service transactions **MUST** be idempotent

---

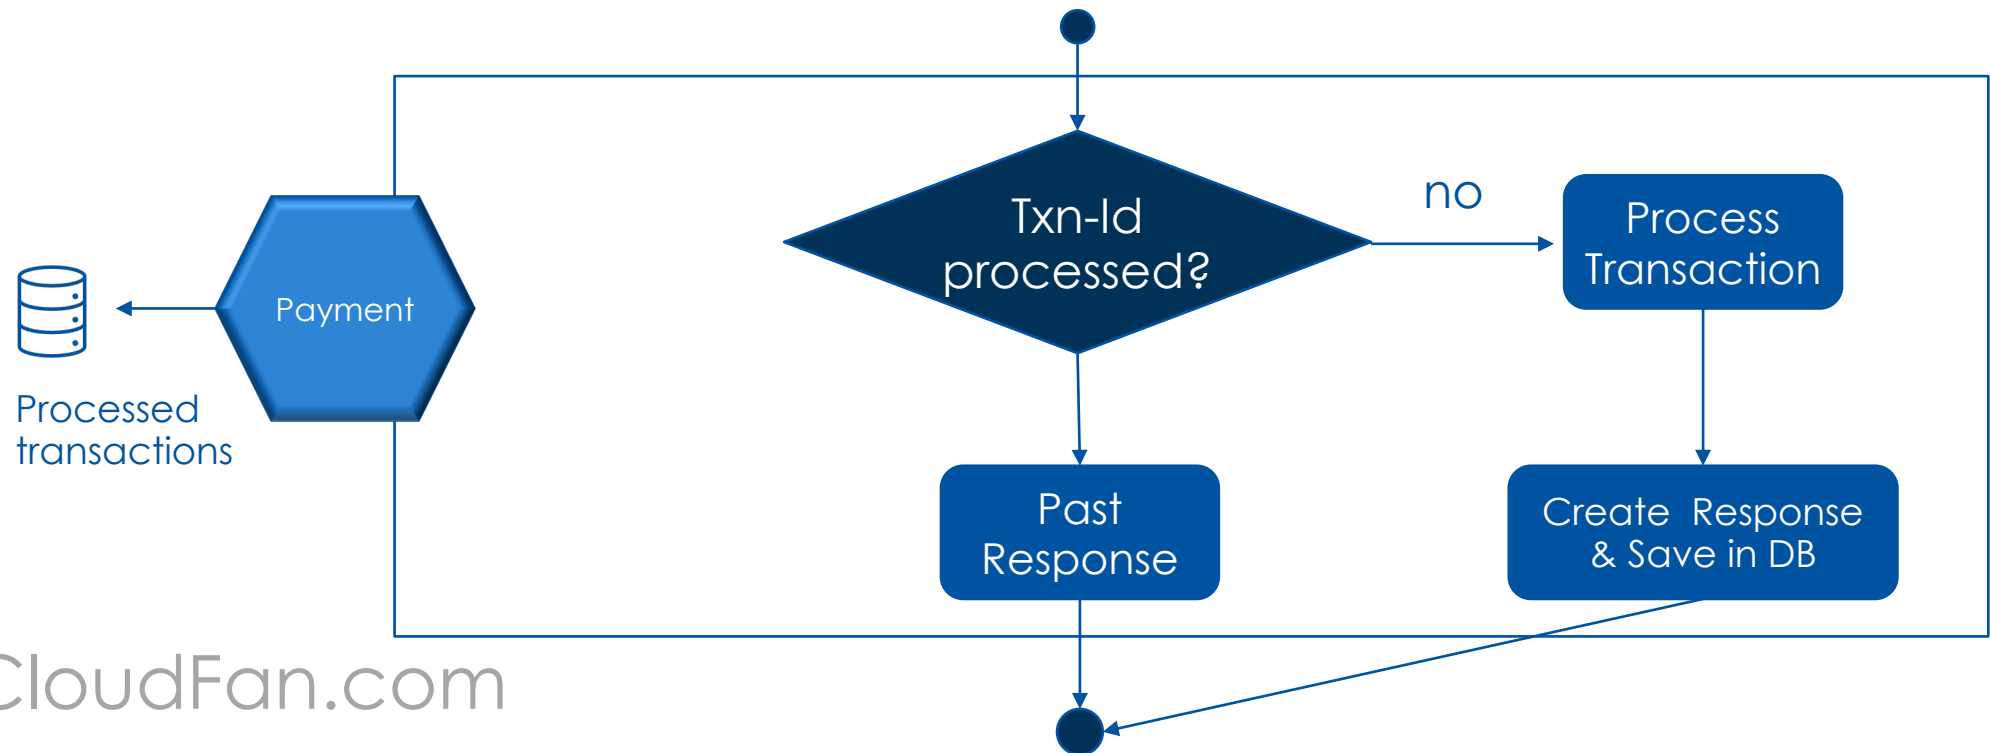
**Idempotence** is the property of certain operations in mathematics and computer science whereby they can be applied multiple times without changing the result beyond the initial application. — [Wikipedia](#)

---

# Idempotent

Services transactions *MUST* be idempotent

- Using a Unique ID for each message can help



## Service Failures

Request failure OK if followed by compensating transactions

- Compensating Transactions CANNOT fail
- Consider event sourcing and state management for services



## Quick Review

### Services

Asynchronous/Messaging is preferred

Each transactions have a unique identity

Operations are idempotent

Compensating transactions cannot fail

# SAGA based booking

Simulation to demonstrate working of SAGA



- 1 Booking SAGA Flow
- 2 Success & Failure scenarios
- 3 Booking SAGA flow in action



IT Lead

## Design the booking process as a SAGA

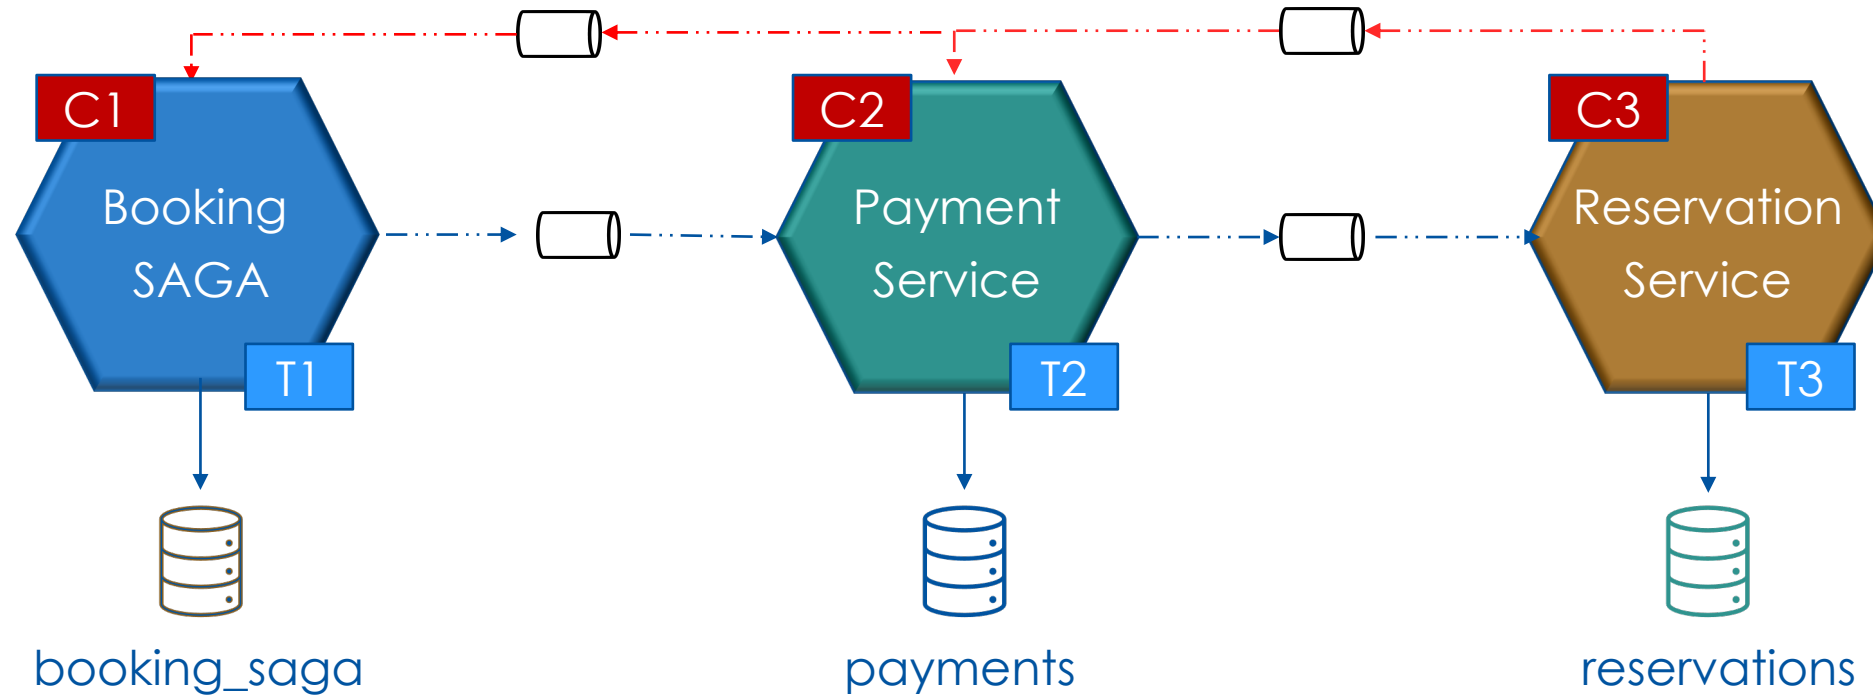
SAGA pattern is difficult to implement without a framework

We will build **version 1** of Booking SAGA without framework

This will help us understand the flow of SAGA



# Booking Flow as a SAGA



T1 add\_booking

T2 process\_payment

T3 process\_reservations

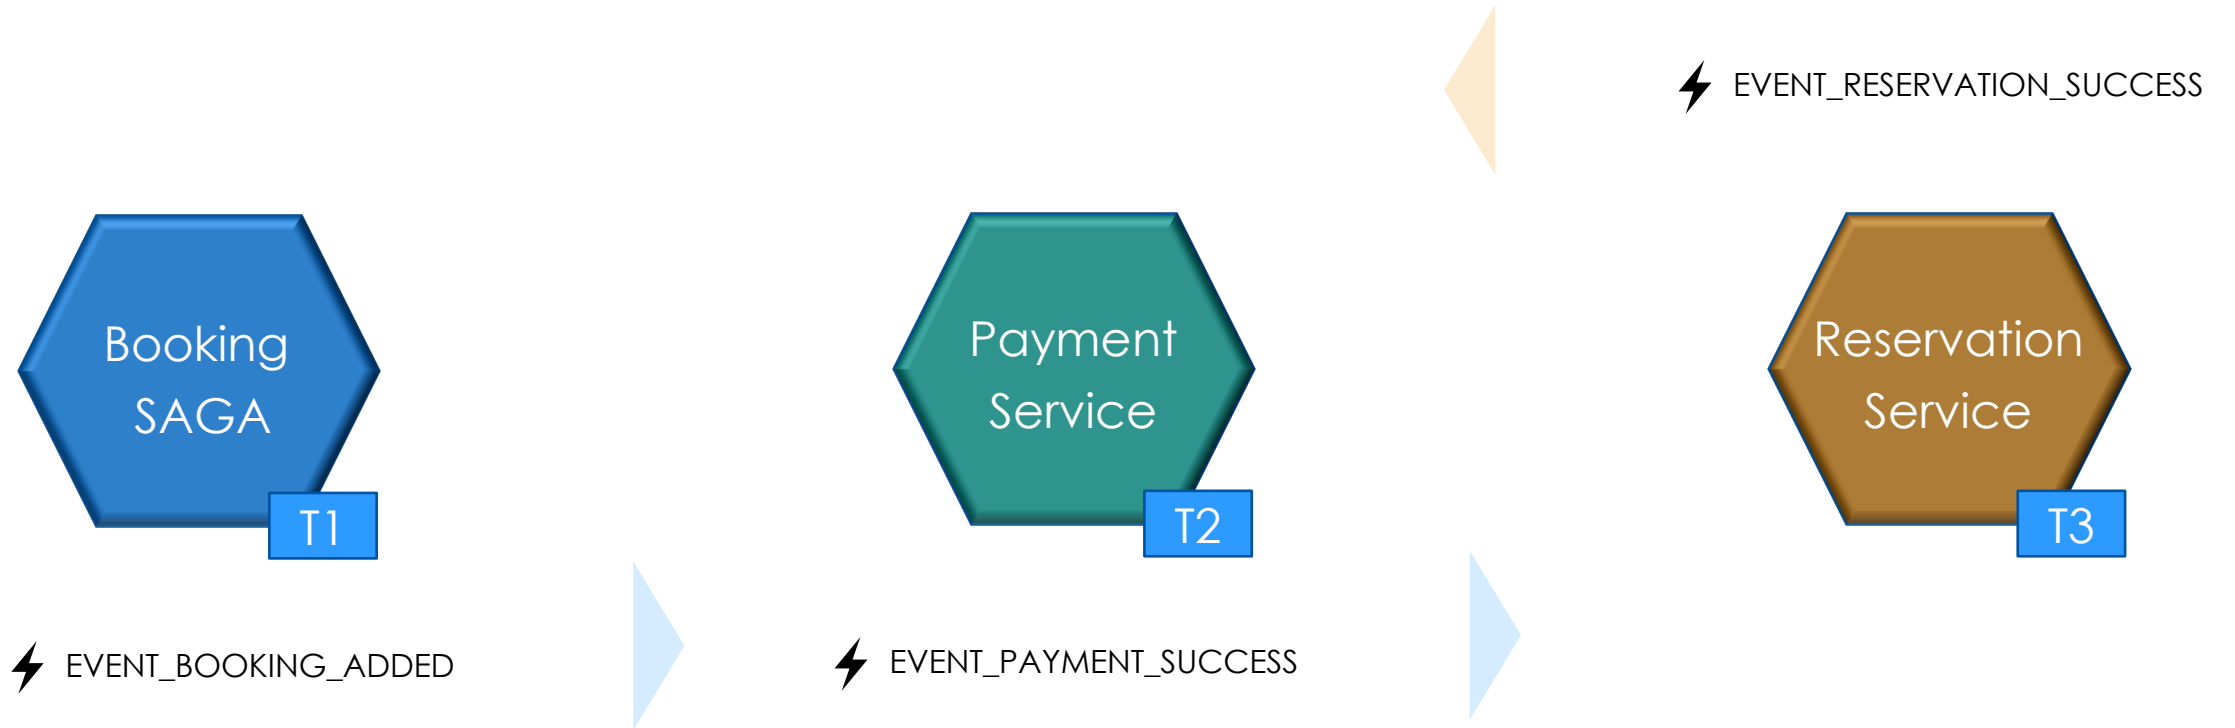
C1 mark\_booking\_failed

C2 refund\_payment

C3 cancel\_reservations

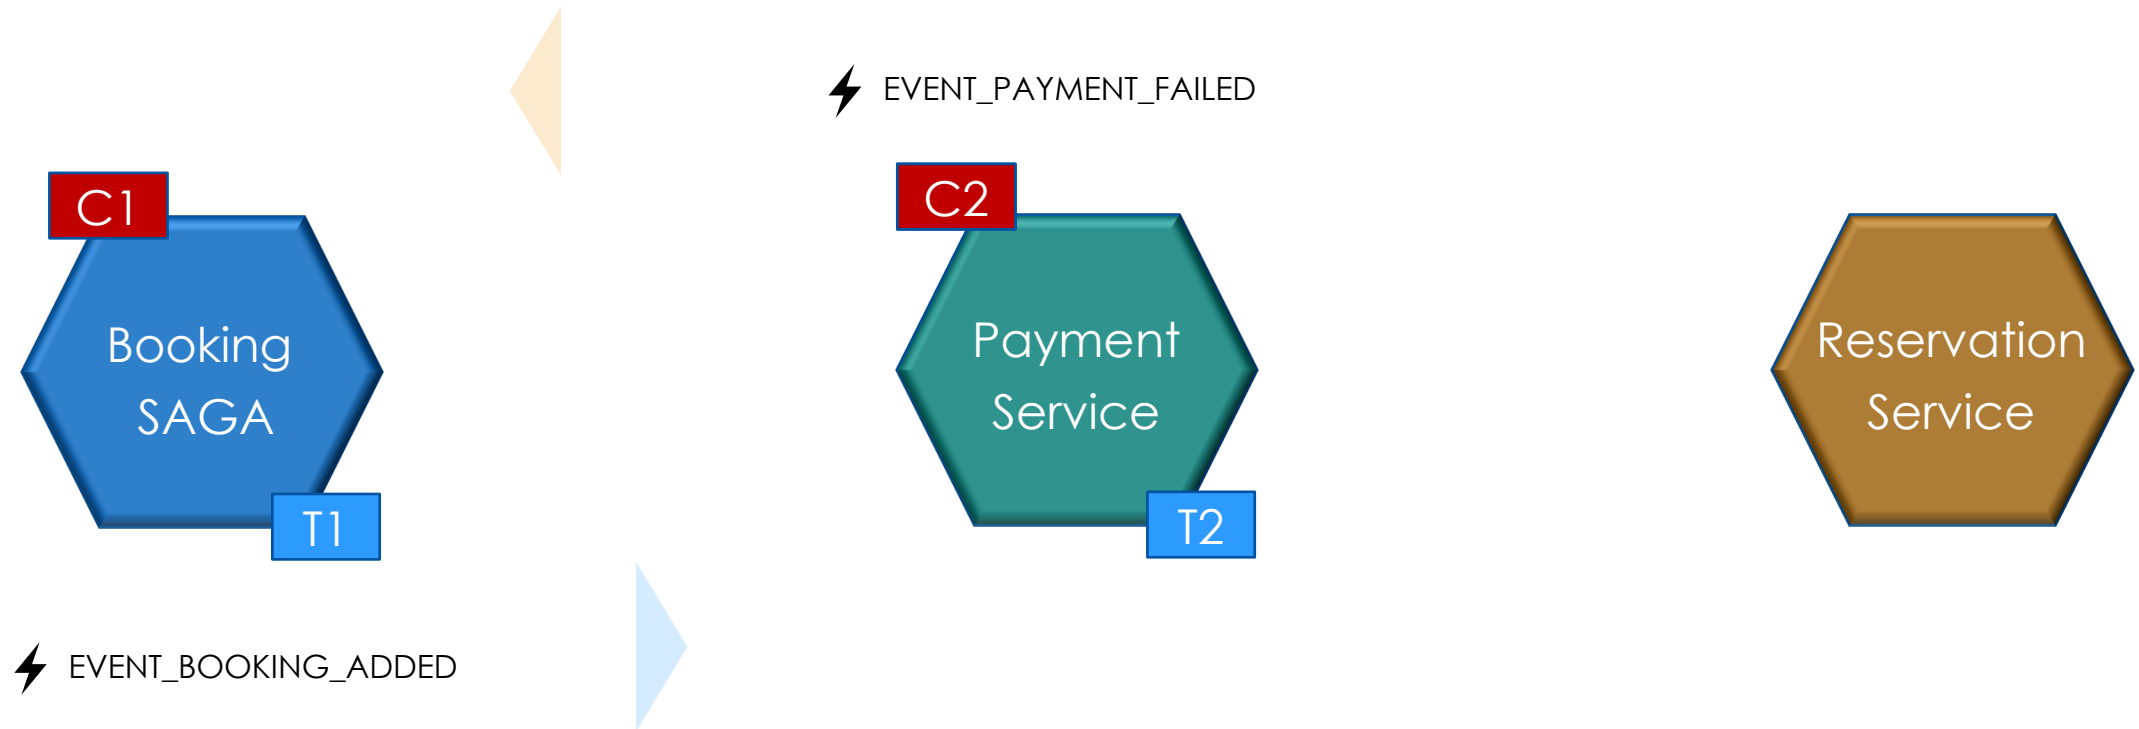
## Events flow

### Successful reservations scenario



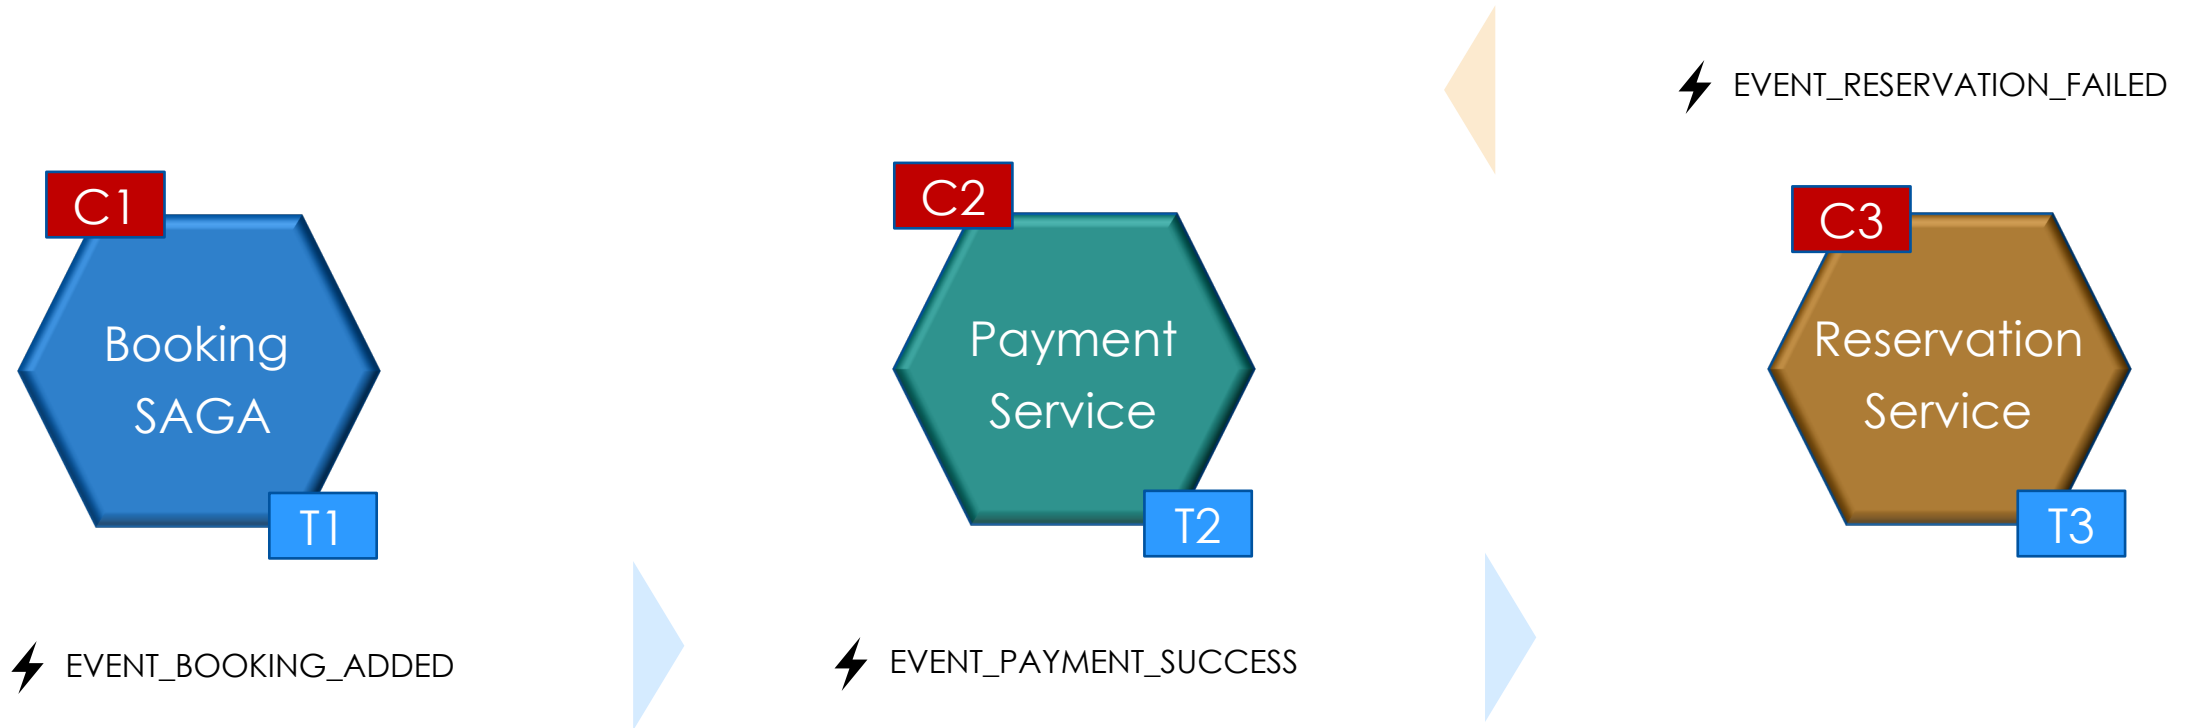
## Events flow

### Payment processing failure scenario



# Events flow

## Reservations failure scenario



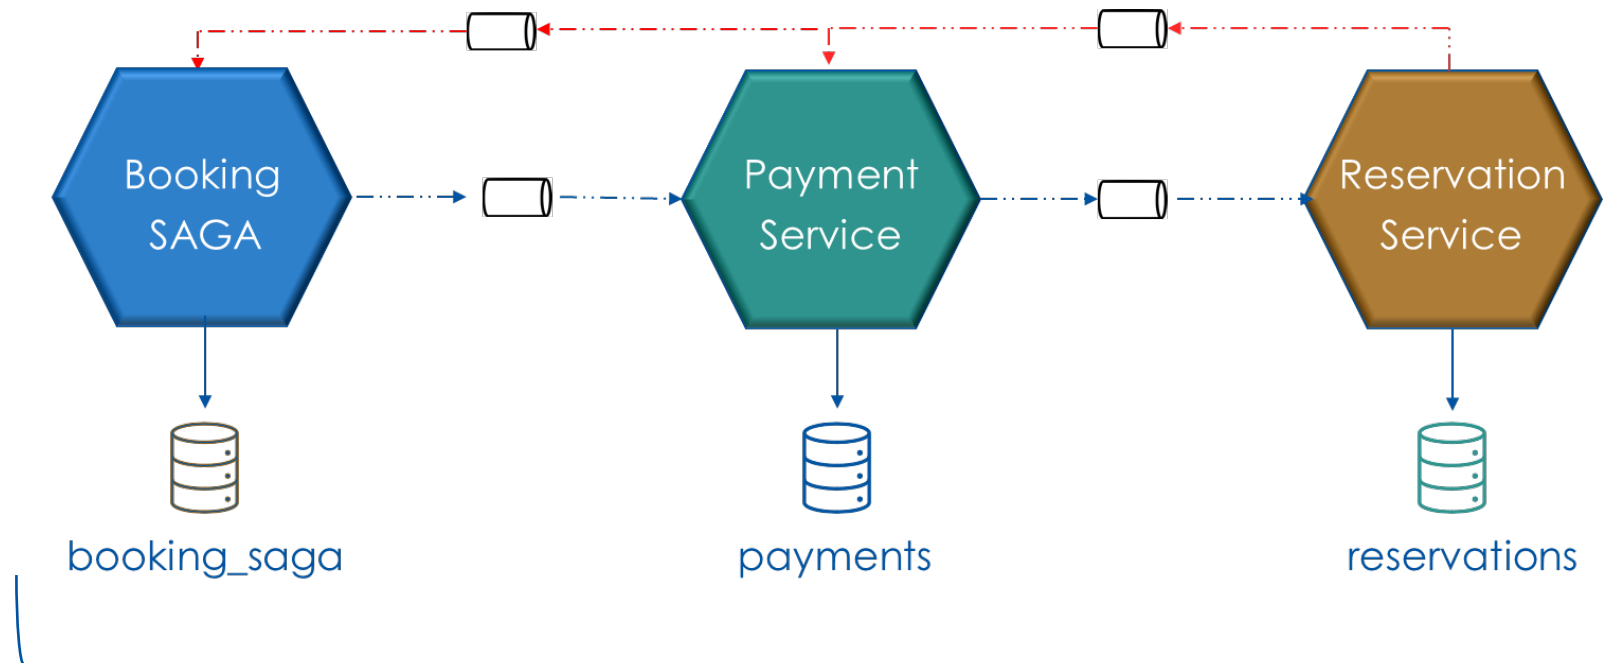
# Implementation stack



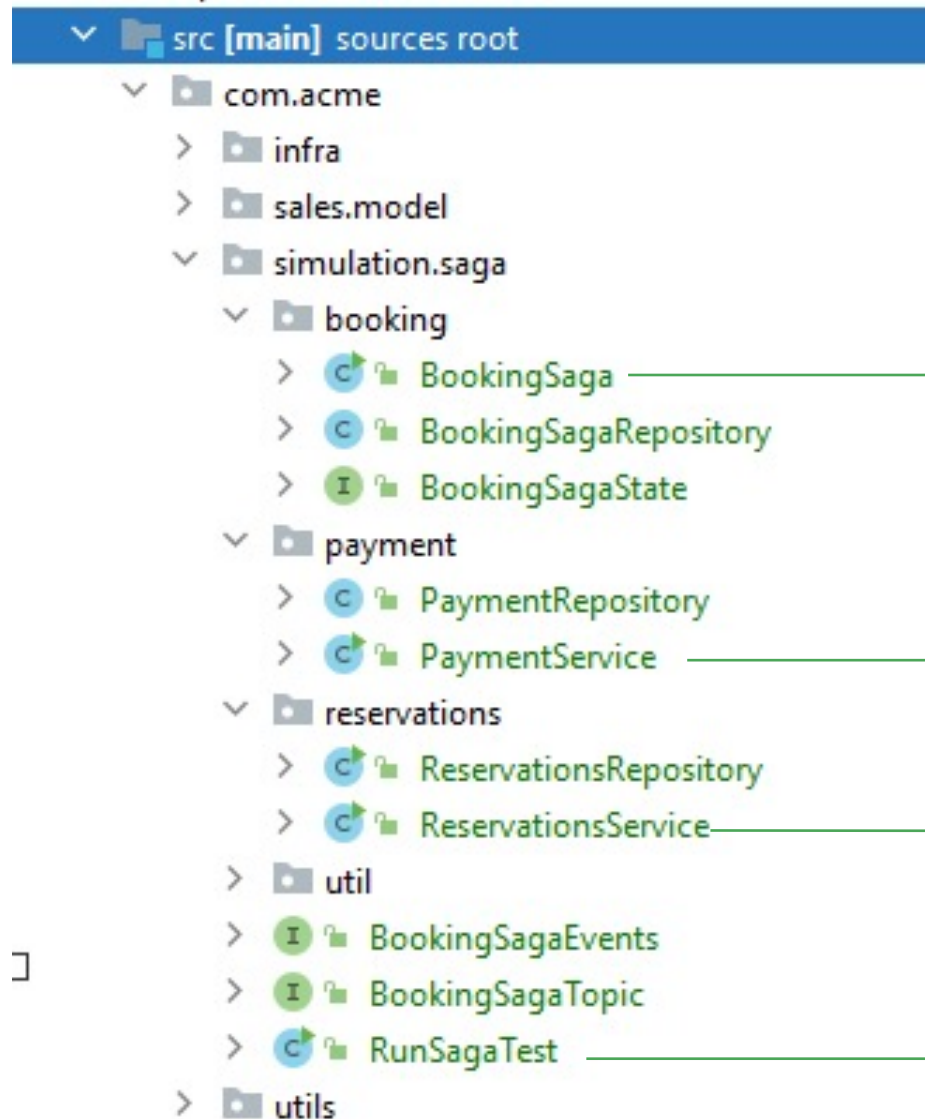
Common topic on Kafka

- Topic `[user name]-bookingsaga`

- Key `booking_id`



# Demo



Manages state of booking SAGA

Simulates processing of payments

Simulates processing of reservations

Starts the Booking SAGA for booking\_id=?

# Demo

src [main] sources root

com.acme

infra

sales.model

simulation.saga

booking

payment

reservations


util

BookingSagaEvents

BookingSagaTopic

RunSagaTest

utils

 cloudkarafka 

Test Instance ▾

CloudKarafka Manager ↗

DETAILS

TOPICS

KAFKA CONFIGURATION

ALARMS

NODES

MFTRICS

Topics

Name	Partitions	Replicas
eeaygrnb-new	2	1
eeaygrnb-bookingsaga	2	1
eeaygrnb-default	5	3

]

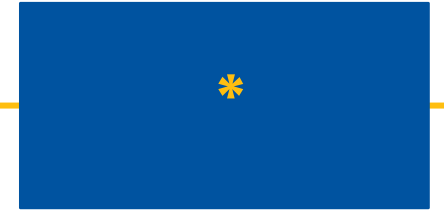
## Next steps

- Design | Code walkthrough of JAVA implementation
- Setup & try out the flow on your local machine



# SAGA Implementation walkthrough

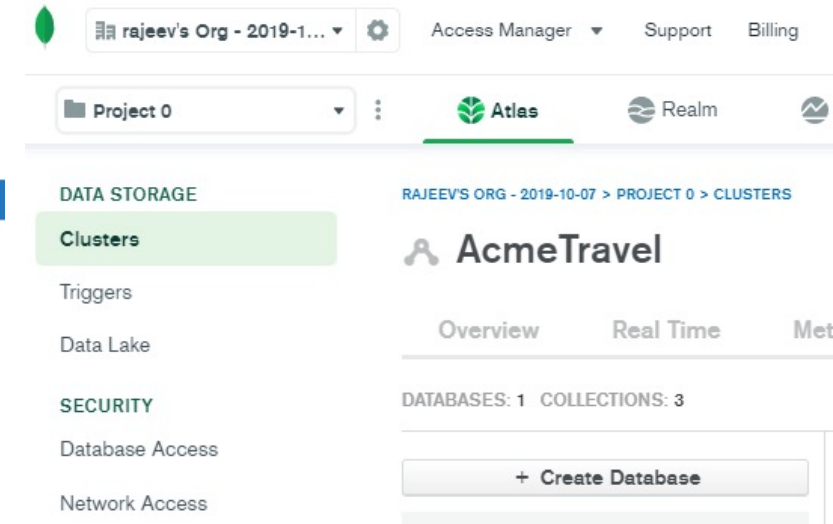
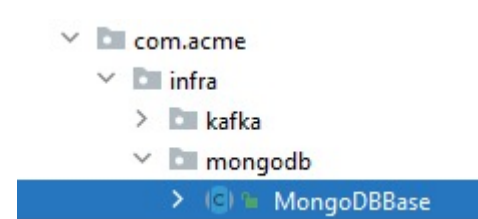
Distributed Booking SAGA implementation



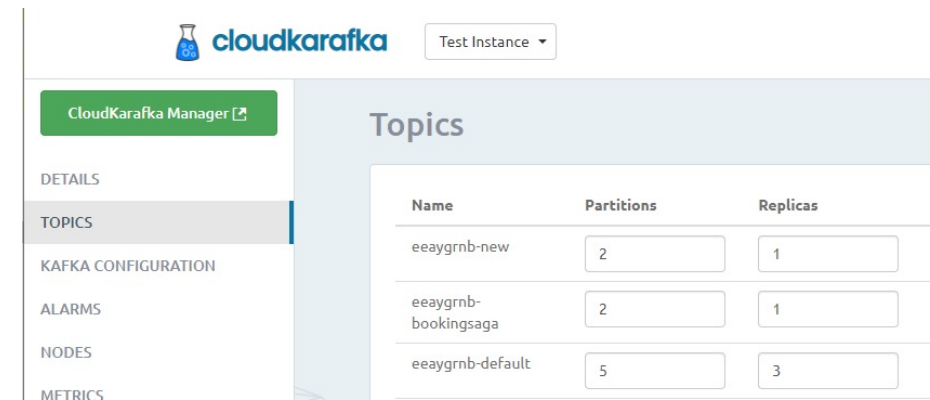
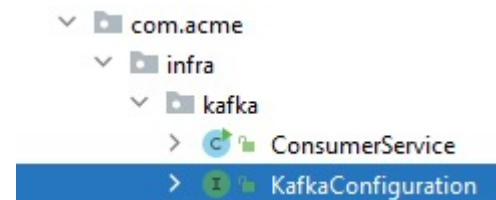
- 1 Walkthrough of Class & State diagrams
- 2 Quick code walkthrough
- 3 SAGA testing steps

# Pre-Requisites

- MongoDB setup is in place



- Cloud Karafka setup is in place

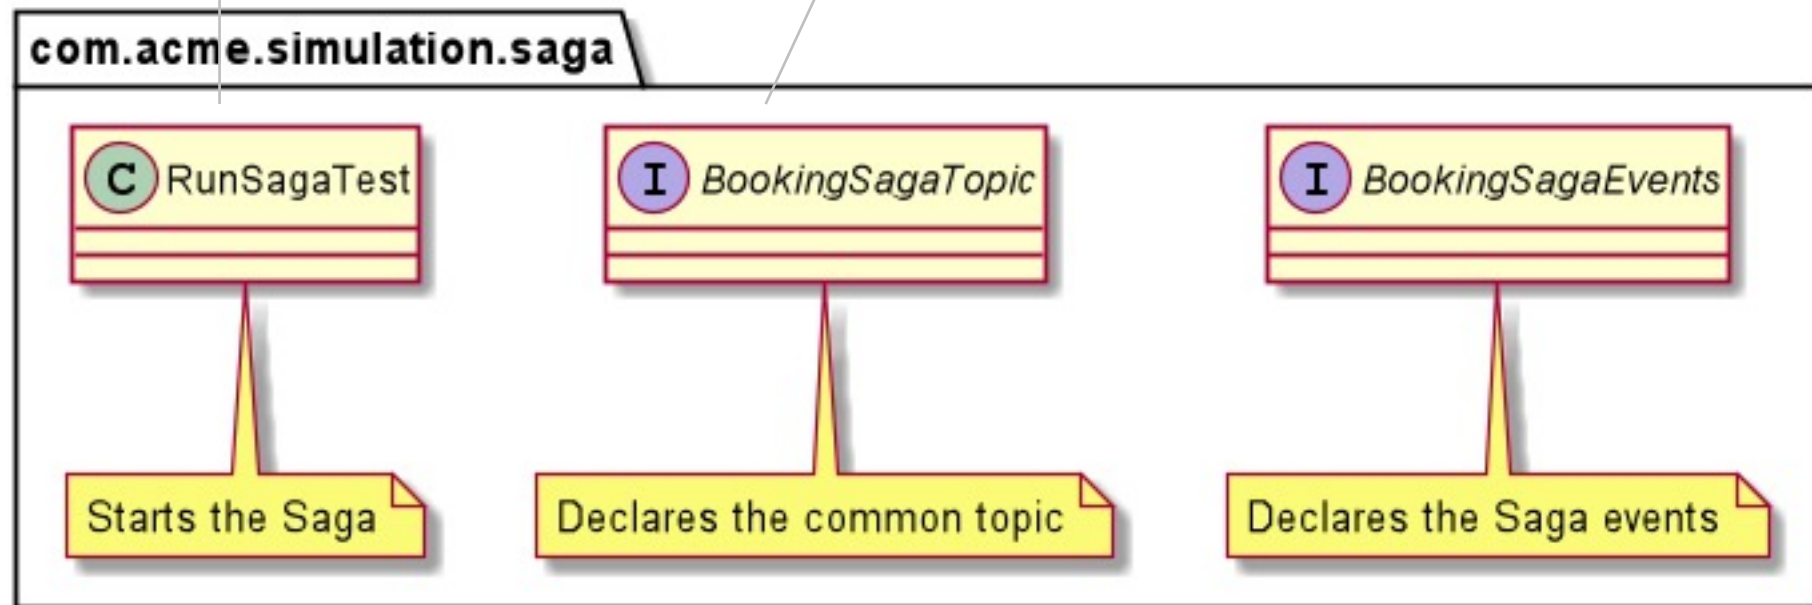


# Class Diagram

/uml/saga/booking.test.classes.puml

- Execute this class for testing; MUST use different booking\_id each time

- Set the topic name in this interface

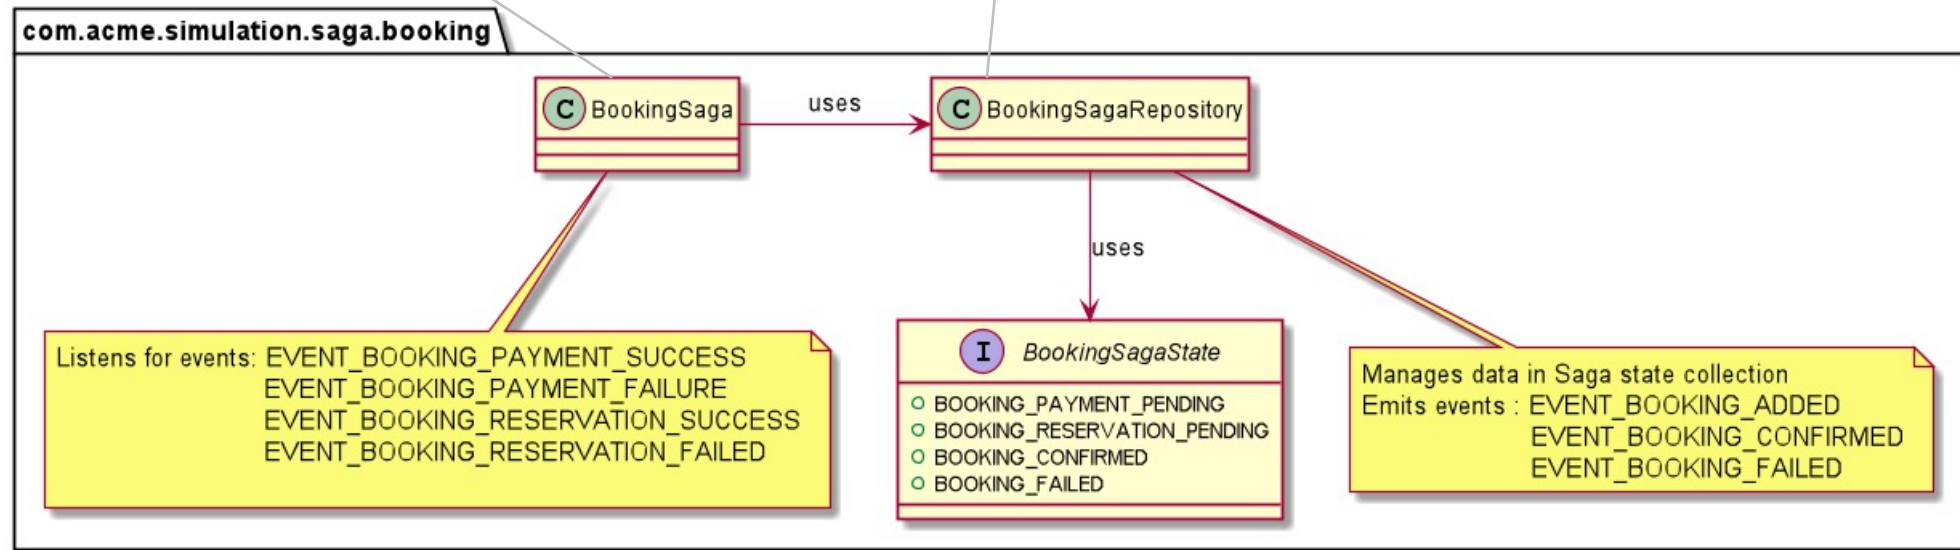


# Class Diagram

/uml/saga/booking.simuation.classes.puml

- Handles the events

- Repository extends *MongoDBBase*
- Events are published from the Repository

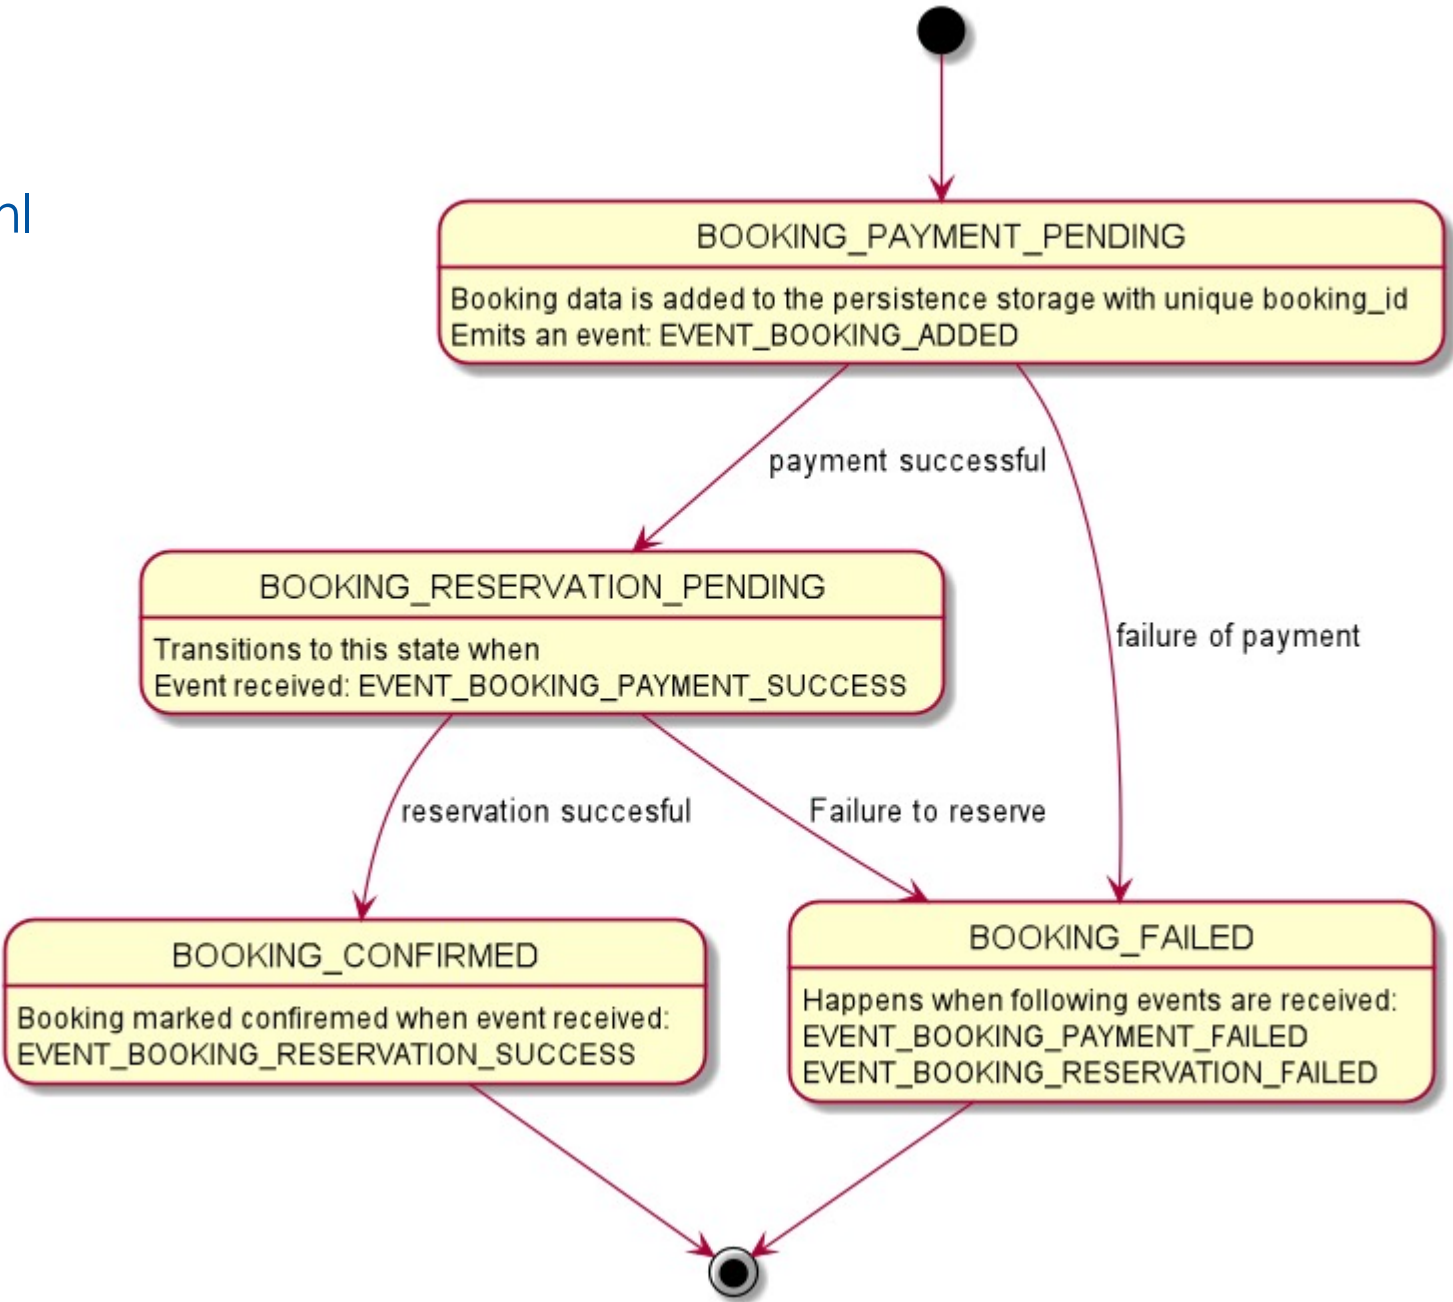


# State Diagram

/uml/saga/booking.saga.state.puml

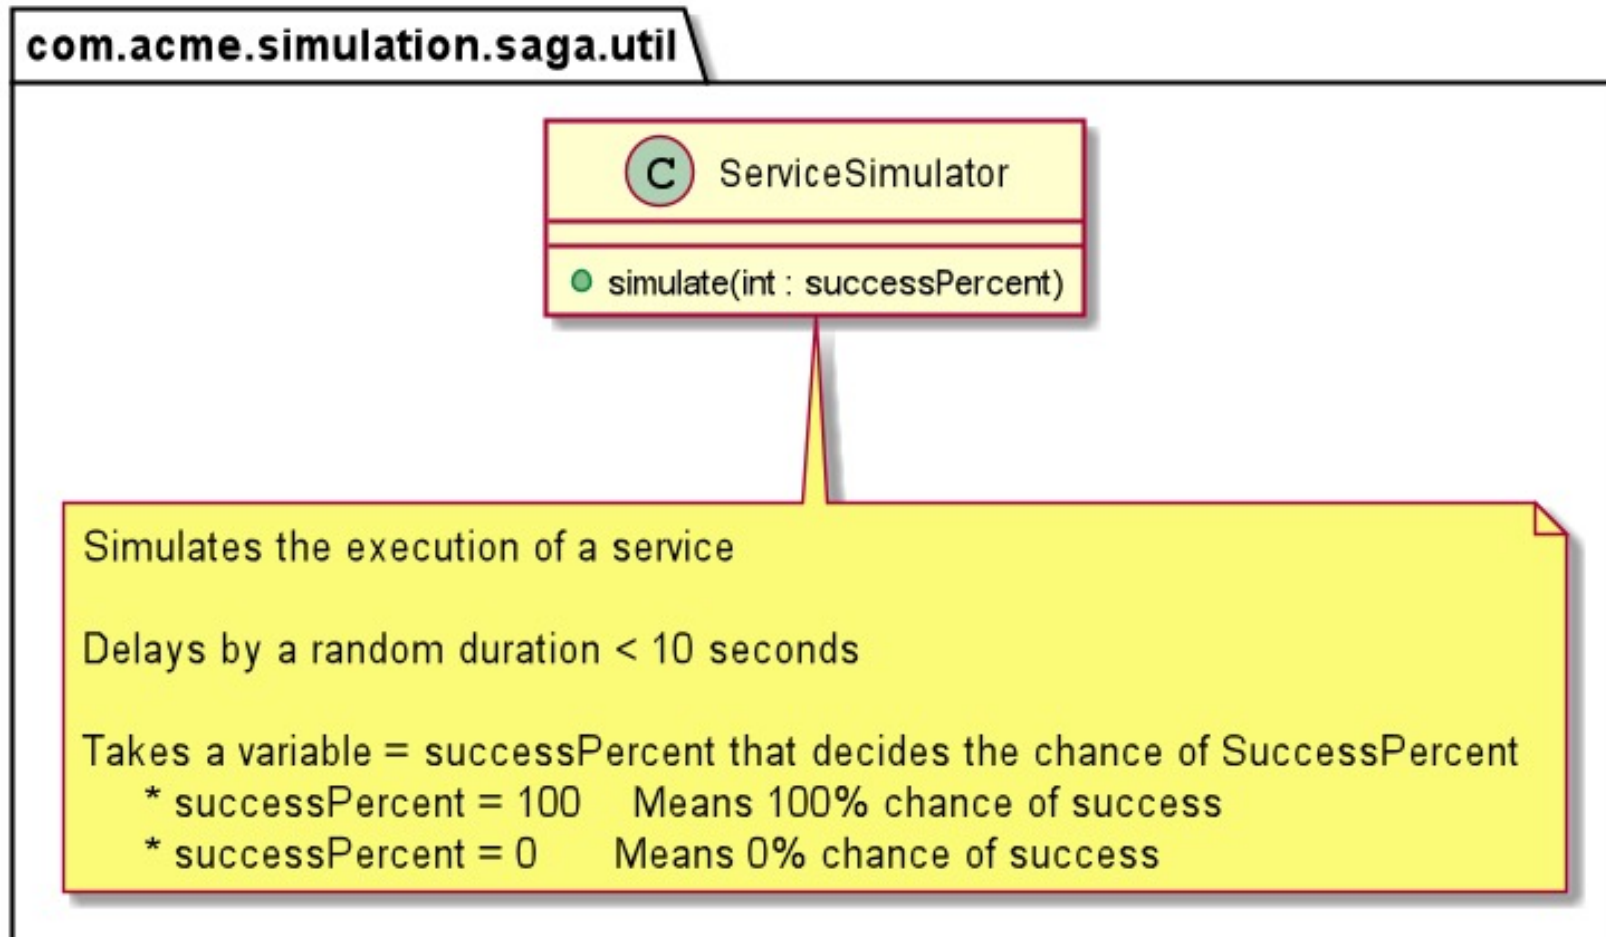


SAGA Execution Coordinator



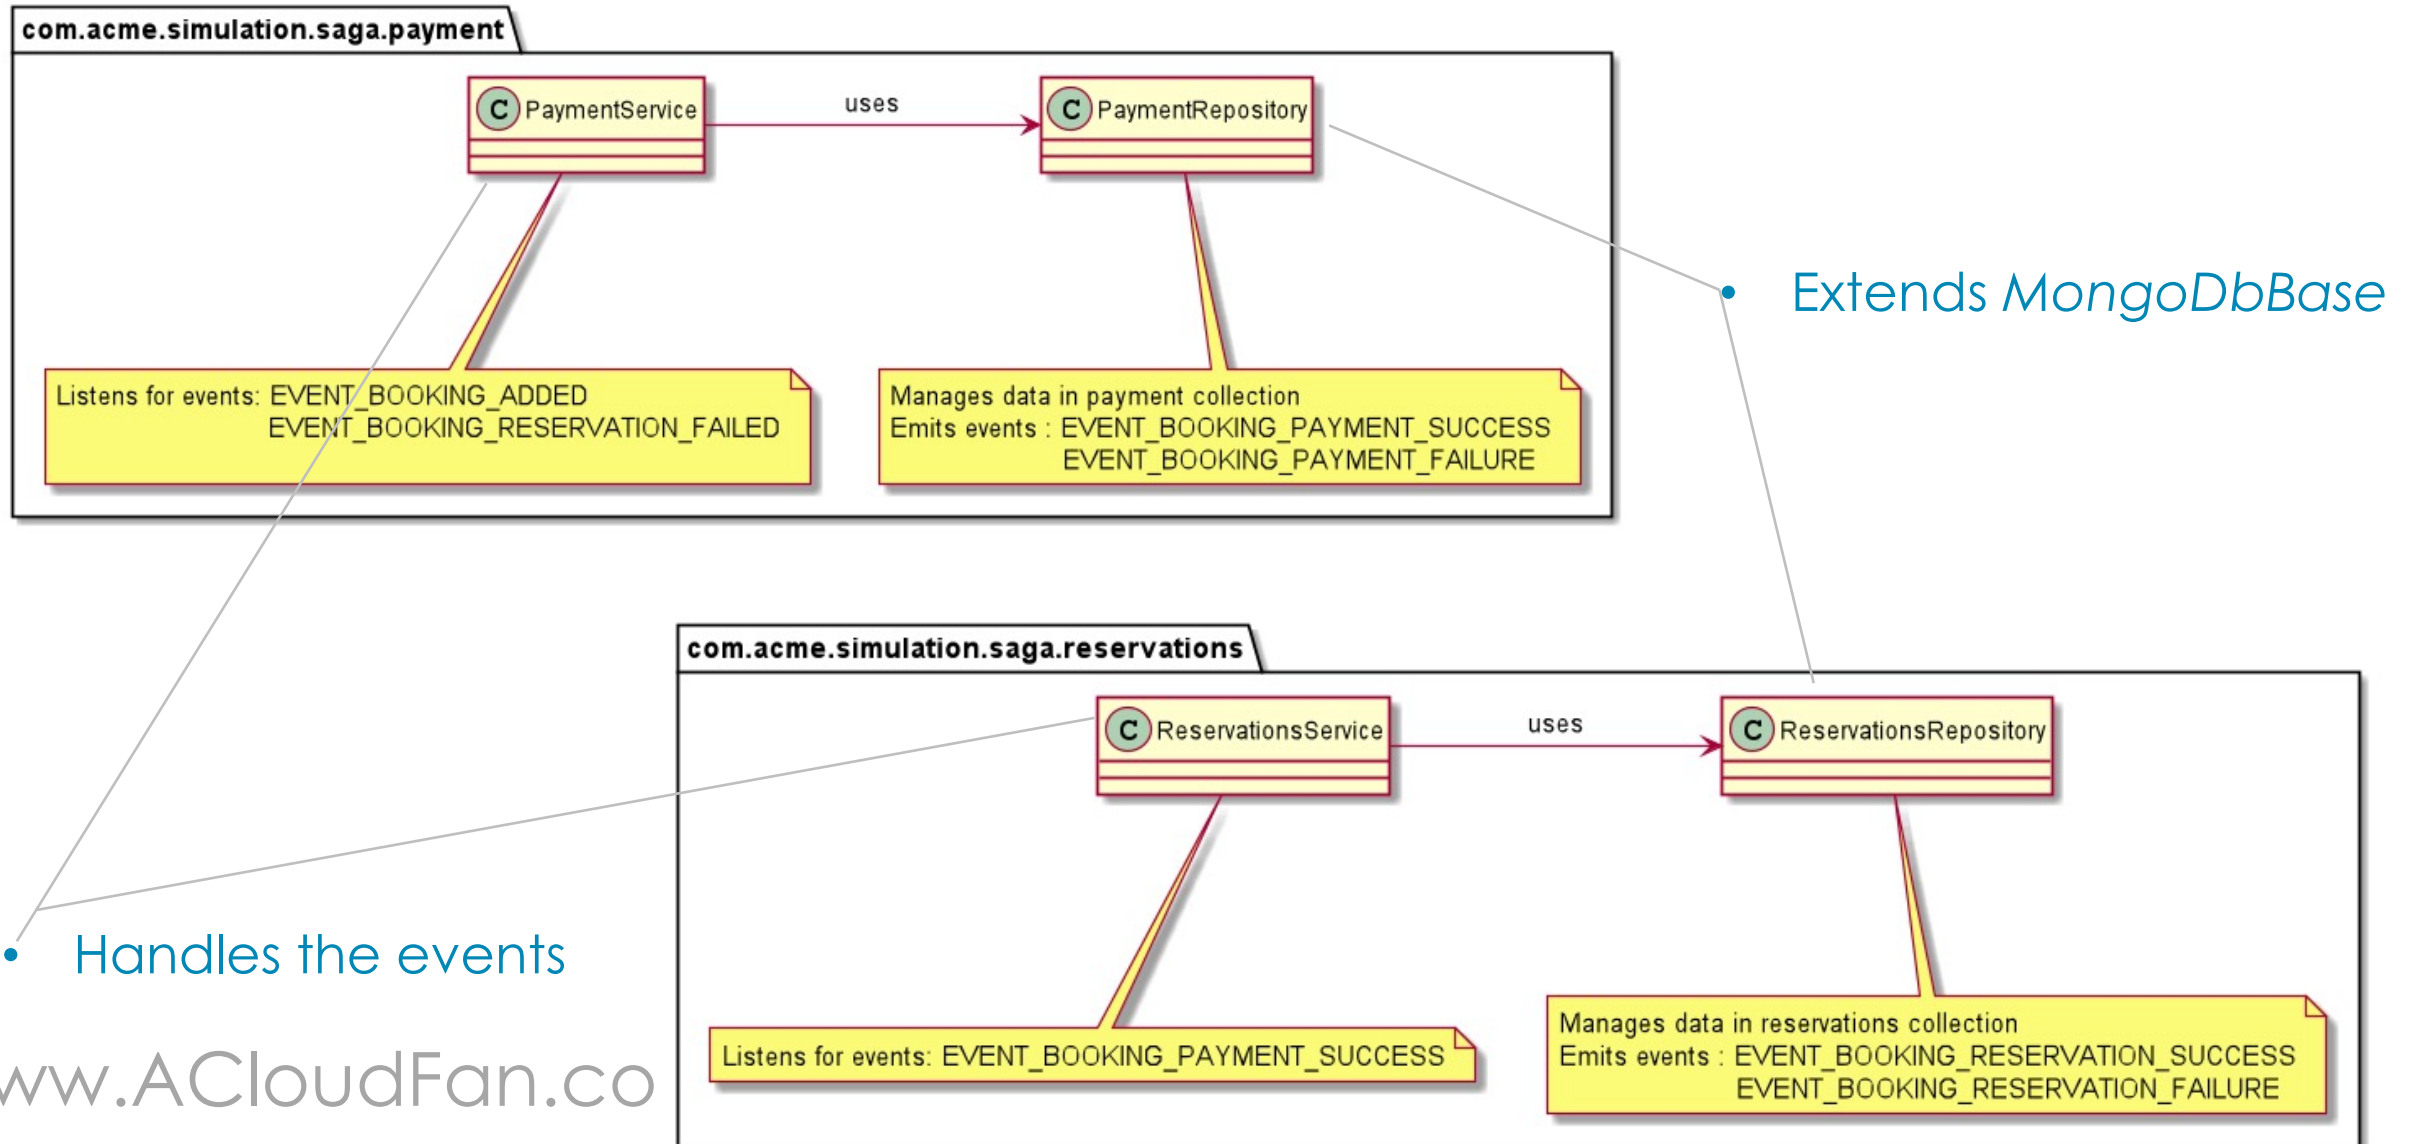
# Service simulation class

/uml/saga/simulator.class.puml



# Class Diagram

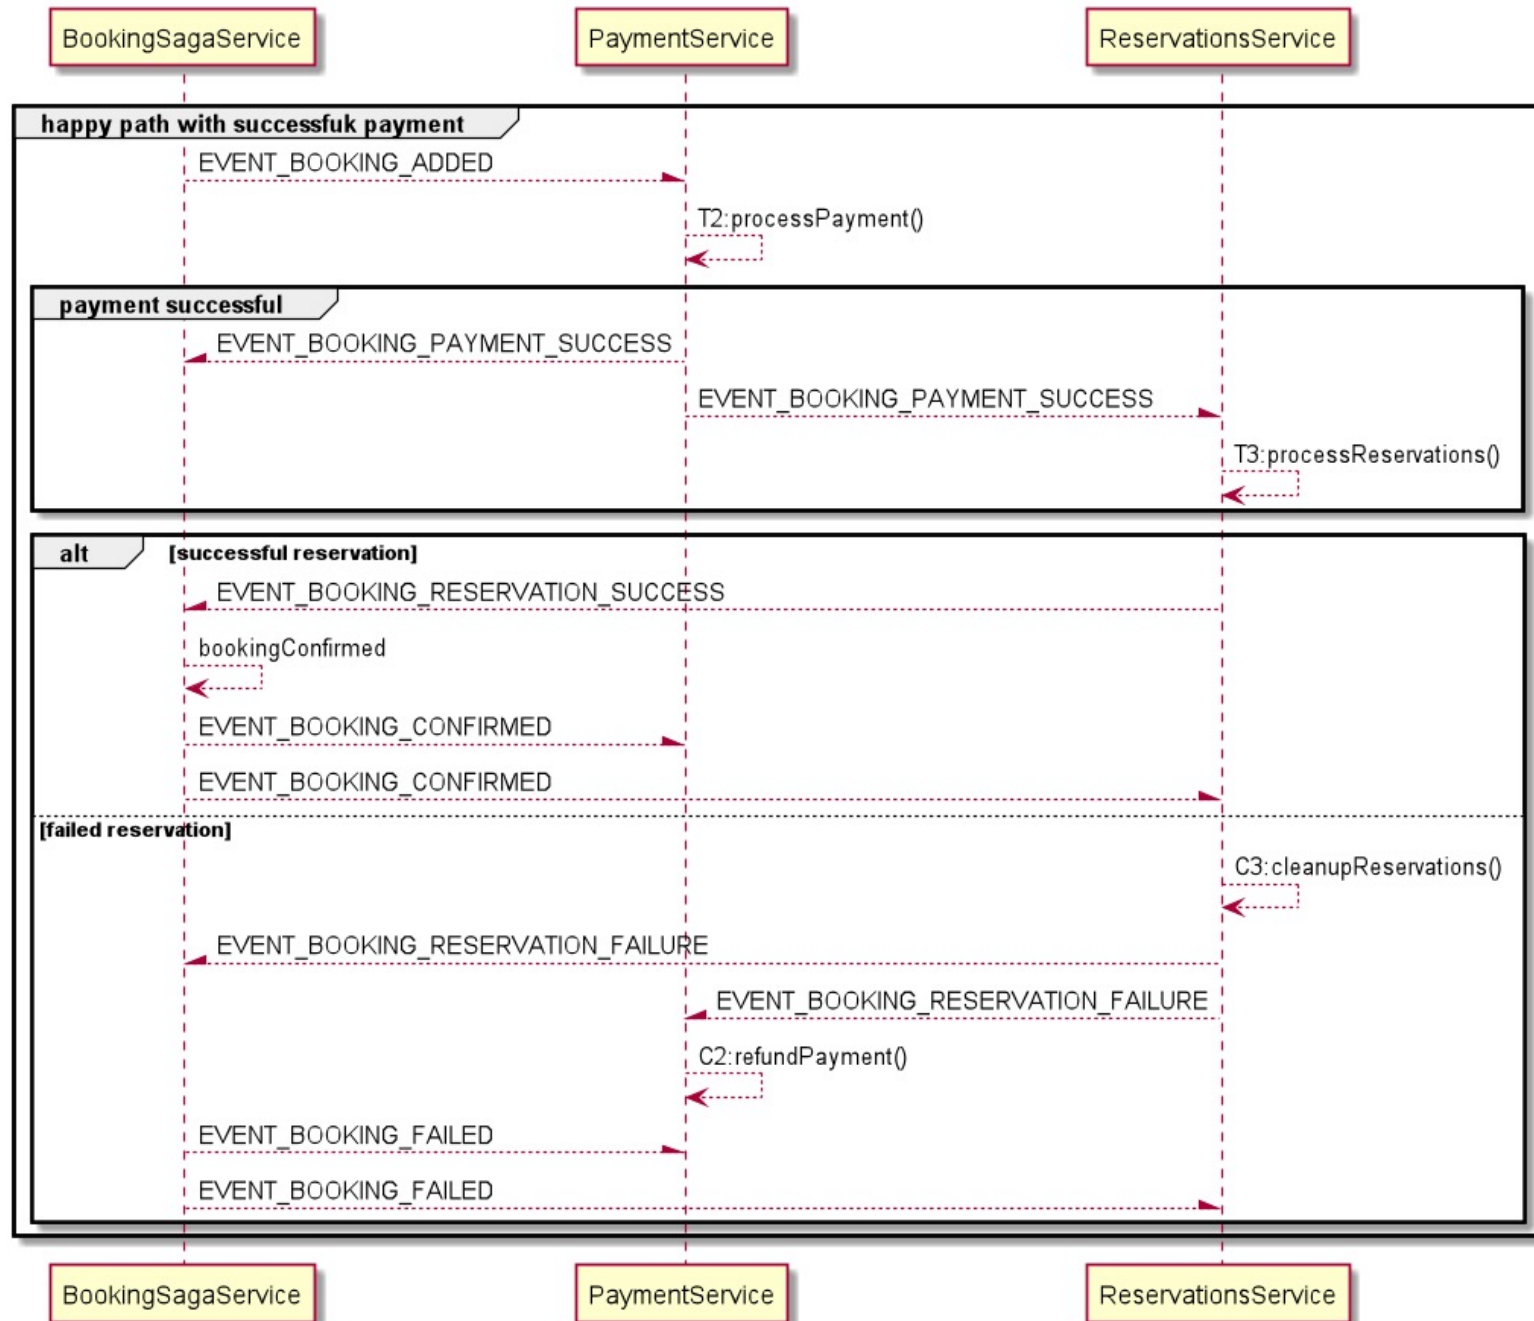
/uml/saga/booking.simuation.classes.puml





# Sequence

/uml/saga/booking.event.sequence.puml

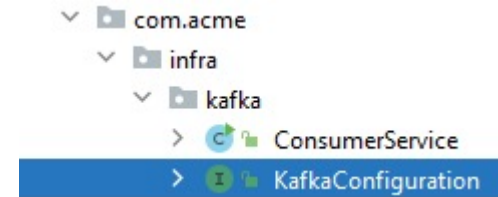




# Test Steps

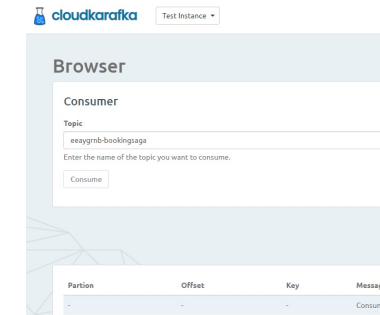
1

Setup kafka topic : booksaga



2

Use Kafka consumer to observe messages



3

Launch the BookingSaga, PaymentService & ReservationsService

4

Execute RunSagaTest with different booking\_id